

Case study: asphalt

The asphalt data

- 31 asphalt pavements prepared under different conditions. How does quality of pavement depend on these?
- Variables:
 - ▶ pct.a.surf Percentage of asphalt in surface layer
 - ▶ pct.a.base Percentage of asphalt in base layer
 - ▶ fines Percentage of fines in surface layer
 - ▶ voids Percentage of voids in surface layer
 - ▶ rut.depth Change in rut depth per million vehicle passes
 - ▶ viscosity Viscosity of asphalt
 - ▶ run 2 data collection periods: 1 for run 1, 0 for run 2.
- rut.depth response. Depends on other variables, how?

Packages for this section

```
library(MASS, exclude = "select")
library(tidyverse)
library(broom)
library(leaps)
```

Make sure to load MASS before tidyverse (for annoying technical reasons), or to load MASS excluding its select (as above).

Getting set up

```
my_url <- "http://ritsokiguess.site/datafiles/asphalt.txt"  
asphalt <- read_delim(my_url, " ")
```

- Quantitative variables with one response: multiple regression.
- Some issues here that don't come up in "simple" regression; handle as we go. (STAB27/STAC67 ideas.)

The data (some)

asphalt

```
# A tibble: 31 x 7
  pct.a.surf pct.a.base fines voids rut.depth viscosity run
  <dbl>      <dbl>   <dbl> <dbl>      <dbl>      <dbl> <dbl>
1     4.68      4.87    8.4  4.92      6.75      2.8   1
2     5.19      4.5     6.5  4.56     13        1.4   1
3     4.82      4.73    7.9  5.32     14.8      1.4   1
4     4.85      4.76    8.3  4.86     12.6      3.3   1
5     4.86      4.95    8.4  3.78     8.25      1.7   1
6     5.16      4.45    7.4  4.40     10.7      2.9   1
7     4.82      5.05    6.8  4.87     7.28      3.7   1
8     4.86      4.7     8.6  4.83     12.7      1.7   1
9     4.78      4.84    6.7  4.86     12.6      0.92  1
10    5.16      4.76    7.7  4.03     20.6      0.68  1
# i 21 more rows
```

Plotting response “rut depth” against everything else

Same idea as for plotting separate predictions on one plot:

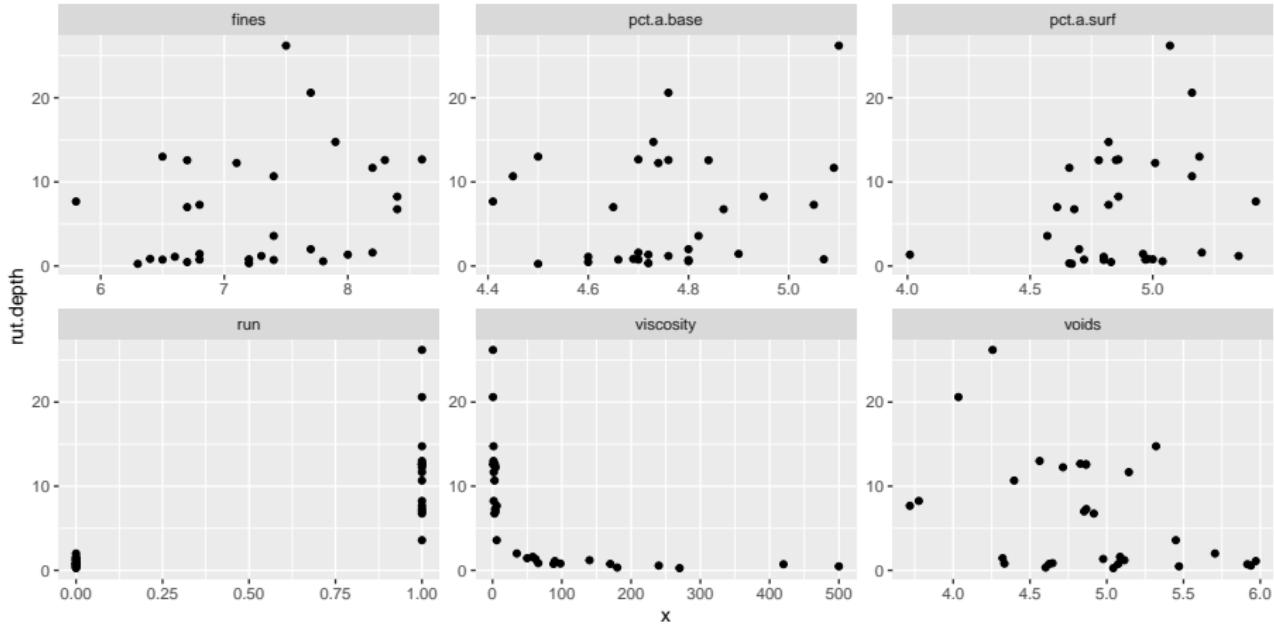
```
asphalt %>%
  pivot_longer(
    -rut.depth,
    names_to = "xname", values_to = "x"
  ) %>%
  ggplot(aes(x = x, y = rut.depth)) + geom_point() +
  facet_wrap(~xname, scales = "free") -> g
```

“collect all the x-variables together into one column called x, with another column xname saying which x they were, then plot these x's against rut.depth, a separate facet for each x-variable.”

I saved this graph to plot later (on the next page).

The plot

g



Interpreting the plots

- One plot of rut depth against each of the six other variables.
- Get rough idea of what's going on.
- Trends mostly weak.
- viscosity has strong but non-linear trend.
- run has effect but variability bigger when run is 1.
- Weak but downward trend for voids.
- Non-linearity of rut.depth-viscosity relationship should concern us.

Log of viscosity: more nearly linear?

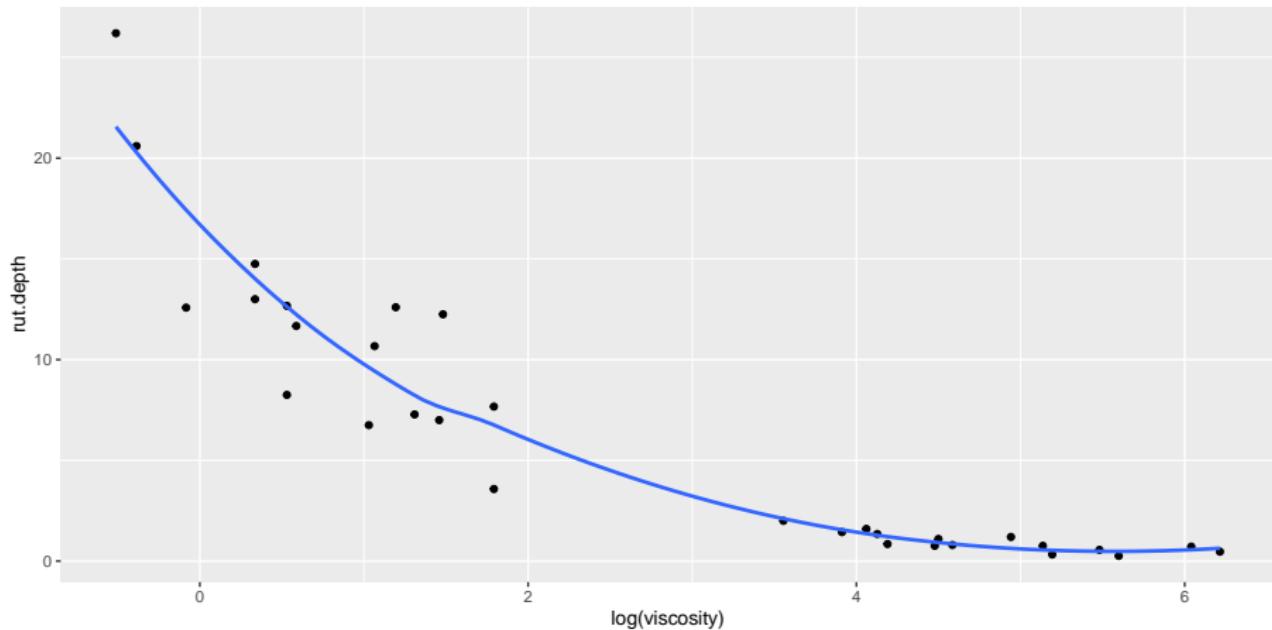
- Take this back to asphalt engineer: suggests log of viscosity:

```
ggplot(asphalt, aes(y = rut.depth, x = log(viscosity))) +  
  geom_point() + geom_smooth(se = FALSE) -> g
```

(plot overleaf)

Rut depth against log-viscosity

g



Comments and next steps

- Not very linear, but better than before.
- In multiple regression, hard to guess which x's affect response. So typically start by predicting from everything else.
- Model formula has response on left, squiggle, explanatorys on right joined by plusses:

```
rut.1 <- lm(rut.depth ~ pct.a.surf + pct.a.base + fines +  
voids + log(viscosity) + run, data = asphalt)
```

Regression output:

```
summary(rut.1)
```

Call:

```
lm(formula = rut.depth ~ pct.a.surf + pct.a.base + fines + voids +
log(viscosity) + run, data = asphalt)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.1211	-1.9075	-0.7175	1.6382	9.5947

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-12.9937	26.2188	-0.496	0.6247
pct.a.surf	3.9706	2.4966	1.590	0.1248
pct.a.base	1.2631	3.9703	0.318	0.7531
fines	0.1164	1.0124	0.115	0.9094
voids	0.5893	1.3244	0.445	0.6604
log(viscosity)	-3.1515	0.9194	-3.428	0.0022 **
run	-1.9655	3.6472	-0.539	0.5949

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.324 on 24 degrees of freedom

Multiple R-squared: 0.806 Adjusted R-squared: 0.7575

Comments

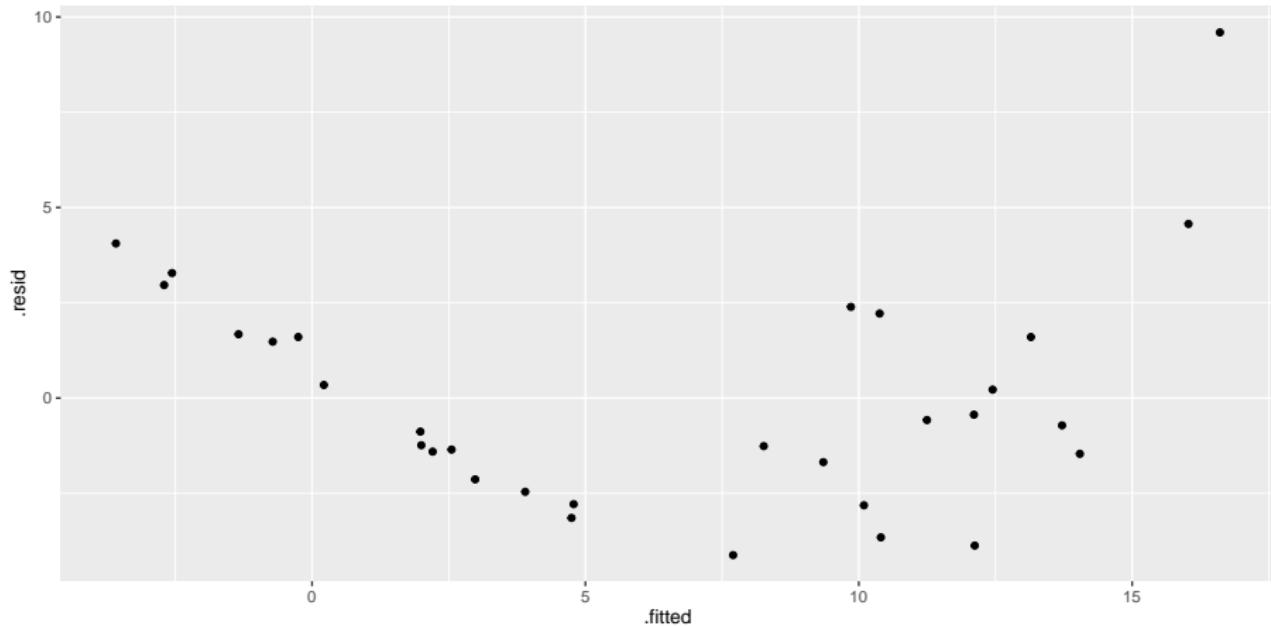
- R-squared 81%, not so bad.
- P-value in `glance` asserts that something helping to predict `rut.depth`.
- Table of coefficients says `log(viscosity)`.
- But confused by clearly non-significant variables: remove those to get clearer picture of what is helpful.

Before we do anything, look at residual plots:

- (a) of residuals against fitted values (as usual)
- (b) of residuals against each explanatory.
- Problem fixes:
 - ▶ with (a): fix response variable;
 - ▶ with some plots in (b): fix those explanatory variables.

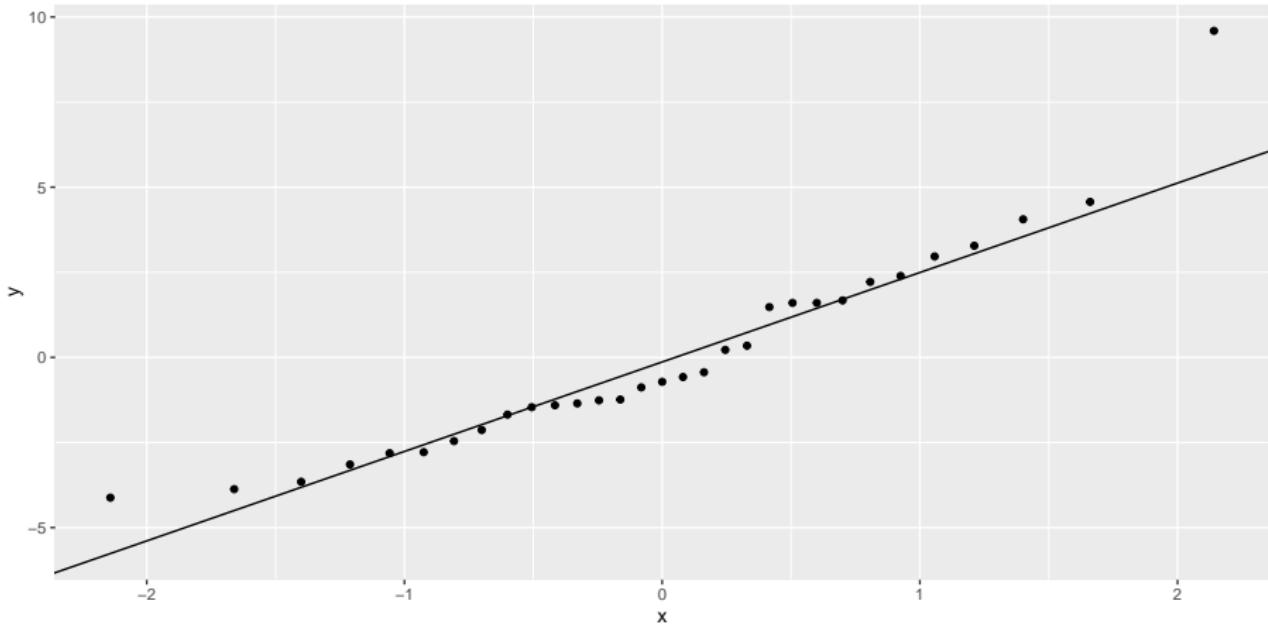
Plot fitted values against residuals

```
ggplot(rut.1, aes(x = .fitted, y = .resid)) + geom_point()
```



Normal quantile plot of residuals

```
ggplot(rut.1, aes(sample = .resid)) + stat_qq() +  
  stat_qq_line()
```



Plotting residuals against x variables

- Problem here is that residuals are in the fitted model, and the observed x -values are in the original data frame asphalt.
- Package broom contains a function augment that combines these two together so that they can later be plotted: start with a model first, and then augment with a data frame:

```
rut.1 %>% augment(asphalt) -> rut.1a  
rut.1a
```

```
# A tibble: 31 x 13  
  pct.a.surf pct.a.base fines voids rut.depth viscosity run .fitted .resid  
    <dbl>      <dbl>   <dbl>  <dbl>     <dbl>      <dbl> <dbl>   <dbl>   <dbl>  
1      4.68      4.87    8.4   4.92      6.75      2.8    1     10.4  -3.65  
2      5.19      4.5     6.5   4.56     13        1.4    1     13.7  -0.718  
3      4.82      4.73    7.9   5.32     14.8       1.4    1     13.1  1.60  
4      4.85      4.76    8.3   4.86     12.6       3.3    1     10.4  2.22  
5      4.86      4.95    8.4   3.78     8.25       1.7    1     12.1  -3.87  
6      5.16      4.45    7.4   4.40     10.7       2.9    1     11.2  -0.577  
7      4.82      5.05    6.8   4.87     7.28       3.7    1     10.1  -2.81  
8      4.86      4.7     8.6   4.83     12.7       1.7    1     12.4  0.221  
9      4.78      4.84    6.7   4.86     12.6       0.92   1     14.0  -1.46  
10     5.16      4.76    7.7   4.03     20.6       0.68   1     16.0  4.57  
# i 21 more rows
```

What does rut.1a contain?

```
names(rut.1a)
```

```
[1] "pct.a.surf"  "pct.a.base"  "fines"      "voids"      "rut.depth"
[6] "viscosity"   "run"        ".fitted"    ".resid"    ".hat"
[11] ".sigma"      ".cooksrd"   ".std.resid"
```

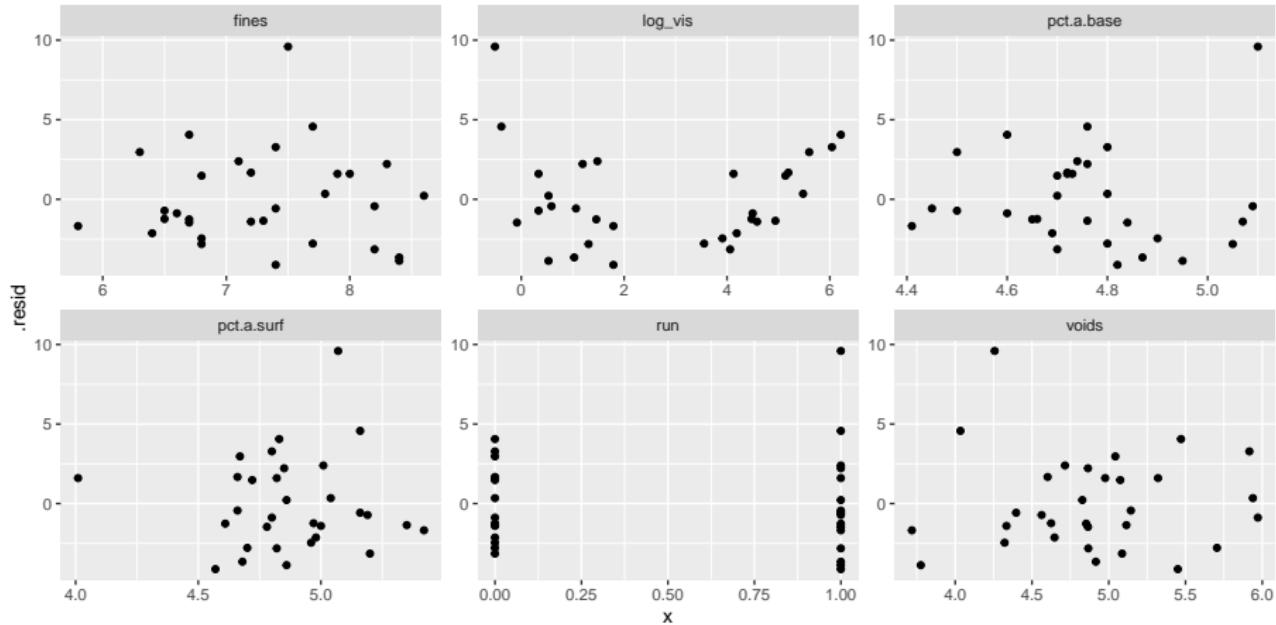
- all the stuff in original data frame, plus:
- quantities from regression (starting with a dot)

Plotting residuals against x -variables

```
rut.1a %>%
  mutate(log_vis=log(viscosity)) %>%
  pivot_longer(
    c(pct.a.surf:voids, run, log_vis),
    names_to="xname", values_to="x"
  ) %>%
  ggplot(aes(x = x, y = .resid)) +
  geom_point() + facet_wrap(~xname, scales = "free") -> g
```

The plot

g



Comments

- There is serious curve in plot of residuals vs. fitted values. Suggests a transformation of y .
- The residuals-vs- x 's plots don't show any serious trends. Worst probably that potential curve against log-viscosity.
- Also, large positive residual, 10, that shows up on all plots. Perhaps transformation of y will help with this too.
- If residual-fitted plot OK, but some residual- x plots not, try transforming those x 's, eg. by adding x^2 to help with curve.

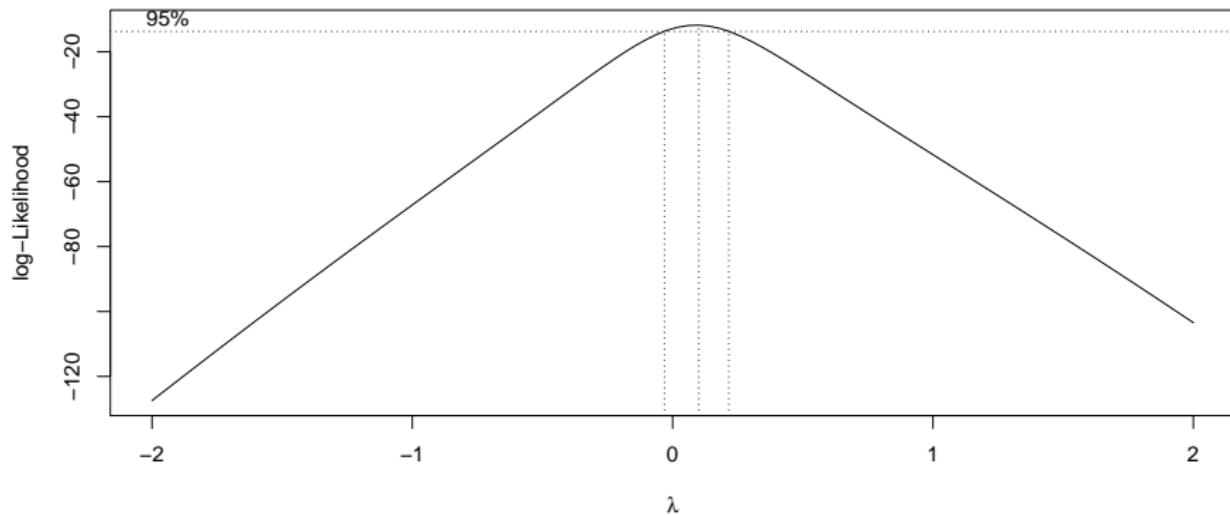
Which transformation?

- Best way: consult with person who brought you the data.
- Can't do that here!
- No idea what transformation would be good.
- Let data choose: “Box-Cox transformation”.
- Scale is that of “ladder of powers”: power transformation, but 0 is log.

Running Box-Cox

From package MASS:

```
boxcox(rut.depth ~ pct.a.surf + pct.a.base + fines + voids +
       log(viscosity) + run, data = asphalt)
```



Comments on Box-Cox plot

- λ represents power to transform y with.
- Best single choice of transformation parameter λ is peak of curve, close to 0.
- Vertical dotted lines give CI for λ , about $(-0.05, 0.2)$.
- $\lambda = 0$ means “log”.
- Narrowness of confidence interval mean that these not supported by data:
 - ▶ No transformation ($\lambda = 1$)
 - ▶ Square root ($\lambda = 0.5$)
 - ▶ Reciprocal ($\lambda = -1$).

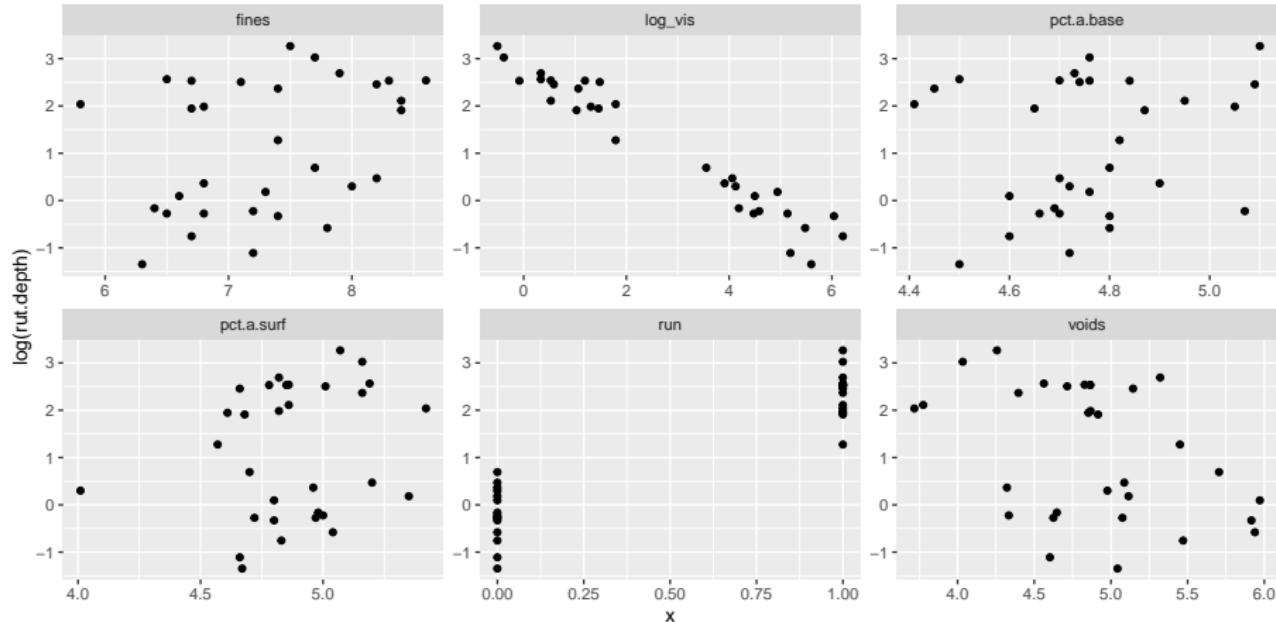
Relationships with explanatories

- As before: plot response (now `log(rut.depth)`) against other explanatory variables, all in one shot:

```
asphalt %>%
  mutate(log_vis=log(viscosity)) %>%
  pivot_longer(
    c(pct.a.surf:voids, run, log_vis),
    names_to="xname", values_to="x"
  ) %>%
  ggplot(aes(y = log(rut.depth), x = x)) + geom_point() +
  facet_wrap(~xname, scales = "free") -> g3
```

The new plots

g3



Modelling with transformed response

- These trends look pretty straight, especially with log.viscosity.
- Values of log.rut.depth for each run have same spread.
- Other trends weak, but are straight if they exist.
- Start modelling from the beginning again.
- Model log.rut.depth in terms of everything else, see what can be removed:

```
rut.2 <- lm(log(rut.depth) ~ pct.a.surf + pct.a.base +  
  fines + voids + log(viscosity) + run, data = asphalt)
```

- use tidy from broom to display just the coefficients.

Output

```
tidy(rut.2)
```

term	estimate	std.error	statistic	p.value
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)	-1.57	2.44	-0.646	0.525
2 pct.a.surf	0.584	0.232	2.52	0.0190
3 pct.a.base	-0.103	0.369	-0.280	0.782
4 fines	0.0978	0.0941	1.04	0.309
5 voids	0.199	0.123	1.62	0.119
6 log(viscosity)	-0.558	0.0854	-6.53	0.000000945
7 run	0.340	0.339	1.00	0.326

Taking out everything non-significant

- Try: remove everything but pct.a.surf and log.viscosity:

```
rut.3 <- lm(log(rut.depth) ~ pct.a.surf + log(viscosity), data = asphalt)
tidy(rut.3)
```

```
# A tibble: 3 x 5
```

term	estimate	std.error	statistic	p.value
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)	0.900	1.08	0.833	4.12e- 1
2 pct.a.surf	0.391	0.219	1.79	8.46e- 2
3 log(viscosity)	-0.619	0.0271	-22.8	1.27e-19

```
summary(rut.3)
```

Call:

```
lm(formula = log(rut.depth) ~ pct.a.surf + log(viscosity), data = asphalt)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.61938	-0.21361	0.06635	0.14932	0.63012

Check that removing all those variables wasn't too much

```
anova(rut.3, rut.2)
```

Analysis of Variance Table

Model 1: $\log(\text{rut.depth}) \sim \text{pct.a.surf} + \log(\text{viscosity})$

Model 2: $\log(\text{rut.depth}) \sim \text{pct.a.surf} + \text{pct.a.base} + \text{fines} + \text{vo}$
run

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	28	2.8809				
2	24	2.2888	4	0.59216	1.5523	0.2191

- H_0 : two models equally good; H_a : bigger model better.
- Null not rejected here; small model as good as the big one, so prefer simpler smaller model rut.3.

Find the largest P-value by eye:

```
tidy(rut.2)
```

```
# A tibble: 7 x 5
```

term	estimate	std.error	statistic	p.value
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)	-1.57	2.44	-0.646	0.525
2 pct.a.surf	0.584	0.232	2.52	0.0190
3 pct.a.base	-0.103	0.369	-0.280	0.782
4 fines	0.0978	0.0941	1.04	0.309
5 voids	0.199	0.123	1.62	0.119
6 log(viscosity)	-0.558	0.0854	-6.53	0.000000945
7 run	0.340	0.339	1.00	0.326

- Largest P-value is 0.78 for pct.a.base, not significant.
- So remove this first, re-fit and re-assess.
- Or, as over.

Get the computer to find the largest P-value for you

- Output from tidy is itself a data frame, thus:

```
tidy(rut.2) %>% arrange(p.value)
```

#	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	log(viscosity)	-0.558	0.0854	-6.53	0.000000945
2	pct.a.surf	0.584	0.232	2.52	0.0190
3	voids	0.199	0.123	1.62	0.119
4	fines	0.0978	0.0941	1.04	0.309
5	run	0.340	0.339	1.00	0.326
6	(Intercept)	-1.57	2.44	-0.646	0.525
7	pct.a.base	-0.103	0.369	-0.280	0.782

- Largest P-value at the bottom.

Take out pct.a.base

- Copy and paste the lm code and remove what you're removing:

```
rut.4 <- lm(log(rut.depth) ~ pct.a.surf + fines + voids +
              log(viscosity) + run, data = asphalt)
summary(rut.4)
```

Call:

```
lm(formula = log(rut.depth) ~ pct.a.surf + fines + voids + log(visco
    run, data = asphalt)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.51610	-0.18785	-0.02248	0.18364	0.57160

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.07850	1.60665	-1.294	0.2076
pct.a.surf	0.59299	0.22526	2.632	0.0143 *
fines	0.08895	0.08701	1.022	0.3165

“Update”

Another way to do the same thing:

```
rut.4 <- update(rut.2, . ~ . - pct.a.base)
tidy(rut.4) %>% arrange(p.value)
```

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	log(viscosity)	-0.552	0.0818	-6.75	0.000000448
2	pct.a.surf	0.593	0.225	2.63	0.0143
3	voids	0.200	0.121	1.66	0.109
4	(Intercept)	-2.08	1.61	-1.29	0.208
5	run	0.360	0.325	1.11	0.279
6	fines	0.0889	0.0870	1.02	0.316

- Again, fines is the one to go. (Output identical as it should be.)

Take out fines:

```
rut.5 <- update(rut.4, . ~ . - fines)  
summary(rut.5)
```

Call:

```
lm(formula = log(rut.depth) ~ pct.a.surf + voids + log(viscosi  
run, data = asphalt)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.57275	-0.20080	0.01061	0.17711	0.59774

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.25533	1.39147	-0.902	0.3753
pct.a.surf	0.54837	0.22118	2.479	0.0200 *
voids	0.23188	0.11676	1.986	0.0577 .

Take out run:

```
rut.6 <- update(rut.5, . ~ . - run)
summary(rut.6)
```

Call:

```
lm(formula = log(rut.depth) ~ pct.a.surf + voids + log(viscosi
    data = asphalt)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.53548	-0.20181	-0.01702	0.16748	0.54707

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.02079	1.36430	-0.748	0.4608
pct.a.surf	0.55547	0.22044	2.520	0.0180 *
voids	0.24479	0.11560	2.118	0.0436 *

Comments

- Here we stop: pct.a.surf, voids and log.viscosity would all make fit significantly worse if removed. So they stay.
- Different final result from taking things out one at a time (top), than by taking out 4 at once (bottom):

coef(rut.6)

(Intercept)	pct.a.surf	voids	log(viscosity)
-1.0207945	0.5554686	0.2447934	-0.6464911

coef(rut.3)

(Intercept)	pct.a.surf	log(viscosity)
0.9001389	0.3911481	-0.6185628

- Point: Can make difference which way we go.

Comments on variable selection

- Best way to decide which x 's belong: expert knowledge: which of them should be important.
- Best automatic method: what we did, “backward selection”.
- Do not learn about “stepwise regression”! **eg. here**
- R has function `step` that does backward selection, like this:

```
step(rut.2, direction = "backward", test = "F")
```

Gets same answer as we did (by removing least significant x).

- Removing non-significant x 's may remove interesting ones whose P-values happened not to reach 0.05. Consider using less stringent cutoff like 0.20 or even bigger.
- Can also fit all possible regressions, as over (may need to do `install.packages("leaps")` first).

All possible regressions (output over)

Uses package leaps:

```
leaps <- regsubsets(log(rut.depth) ~ pct.a.surf +
                      pct.a.base + fines + voids +
                      log(viscosity) + run,
                      data = asphalt, nbest = 2)
s <- summary(leaps)
with(s, data.frame(rsq, outmat)) -> d
```

The output

```
d %>% rownames_to_column("model") %>% arrange(desc(rsq))
```

	model	rsq	pct.a.surf	pct.a.base	fines	voids	log.viscosity.	run
1	6	(1)	0.9609642	*	*	*	*	*
2	5	(1)	0.9608365	*		*	*	*
3	5	(2)	0.9593265	*	*	*		*
4	4	(1)	0.9591996	*			*	*
5	4	(2)	0.9589206	*		*		*
6	3	(1)	0.9578631	*		*		*
7	3	(2)	0.9534561	*		*		*
8	2	(1)	0.9508647	*				*
9	2	(2)	0.9479541			*		*
10	1	(1)	0.9452562					*
11	1	(2)	0.8624107					*

Comments

- Problem: even adding a worthless x increases R-squared. So try for line where R-squared stops increasing “too much”, eg. top line (just log.viscosity), first 3-variable line (backwards-elimination model). Hard to judge.
- One solution (STAC67): adjusted R-squared, where adding worthless variable makes it go down.
- `data.frame` rather than `tibble` because there are several columns in `outmat`.

All possible regressions, adjusted R-squared

```
with(s, data.frame(adjr2, outmat)) %>%
  rownames_to_column("model") %>%
  arrange(desc(adjr2))
```

	model	adjr2	pct.a.surf	pct.a.base	fines	voids	log.viscosity.	run
1	3	(1)	0.9531812	*		*		*
2	5	(1)	0.9530038	*		*	*	*
3	4	(1)	0.9529226	*		*		*
4	4	(2)	0.9526007	*		*	*	*
5	6	(1)	0.9512052	*	*	*		*
6	5	(2)	0.9511918	*	*	*		*
7	3	(2)	0.9482845	*		*		*
8	2	(1)	0.9473550	*				*
9	2	(2)	0.9442365			*		*
10	1	(1)	0.9433685					*
11	1	(2)	0.8576662					*

Revisiting the best model

- Best model was our `rut.6`:

```
tidy(rut.6)
```

```
# A tibble: 4 x 5
  term            estimate std.error statistic p.value
  <chr>          <dbl>     <dbl>      <dbl>    <dbl>
1 (Intercept)   -1.02      1.36     -0.748  4.61e- 1
2 pct.a.surf    0.555     0.220      2.52   1.80e- 2
3 voids         0.245     0.116      2.12   4.36e- 2
4 log(viscosity) -0.646    0.0288    -22.5  5.29e-19
```

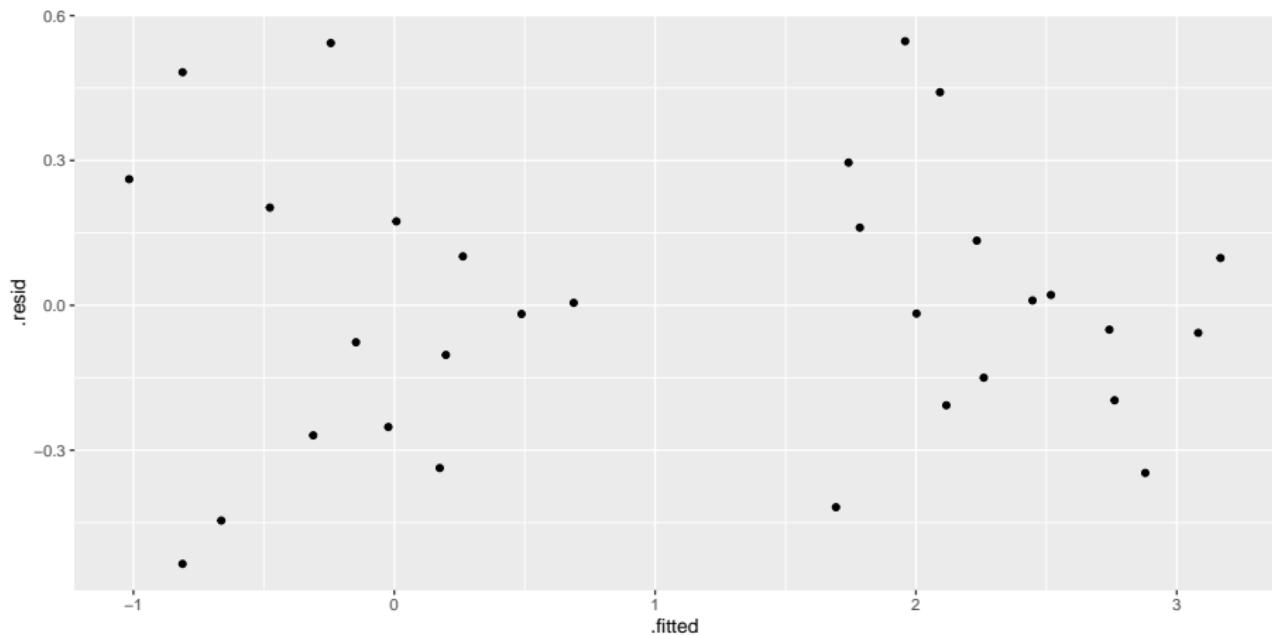
Revisiting (2)

- Regression slopes say that rut depth increases as log-viscosity decreases, pct.a.surf increases and voids increases. This more or less checks out with out scatterplots against log.viscosity.
- We should check residual plots again, though previous scatterplots say it's unlikely that there will be a problem:

```
g <- ggplot(rut.6, aes(y = .resid, x = .fitted)) +  
  geom_point()
```

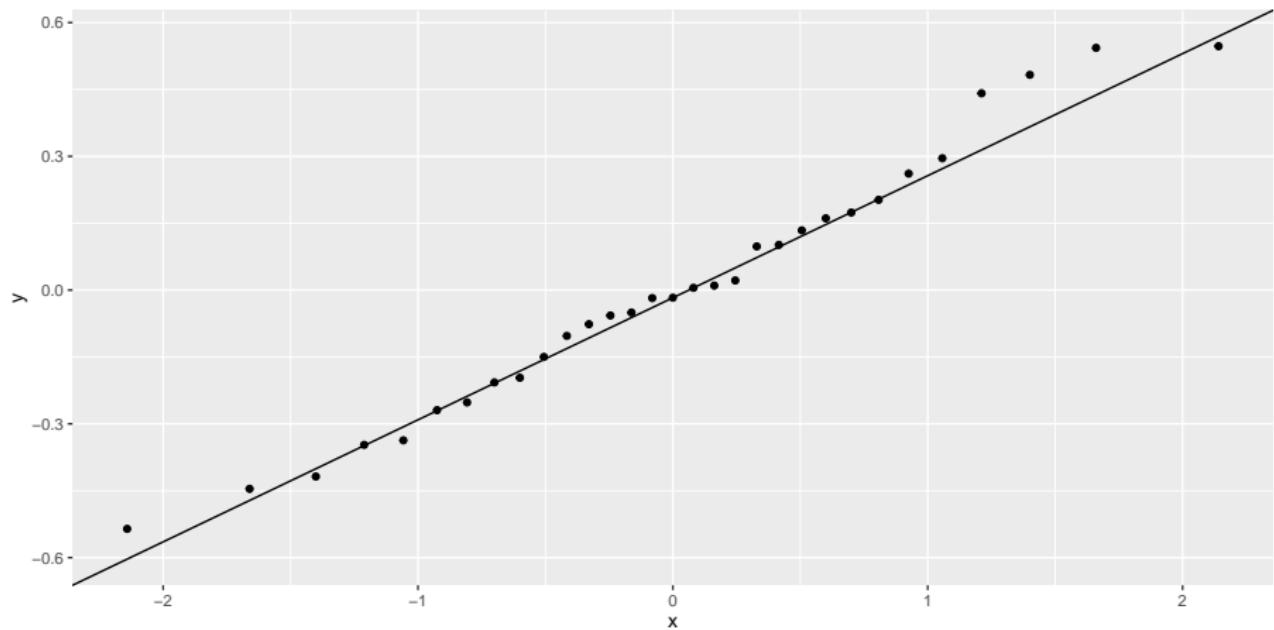
Residuals against fitted values

g



Normal quantile plot of residuals

```
ggplot(rut.6, aes(sample = .resid)) + stat_qq() + stat_qq_line()
```



Plotting residuals against x's

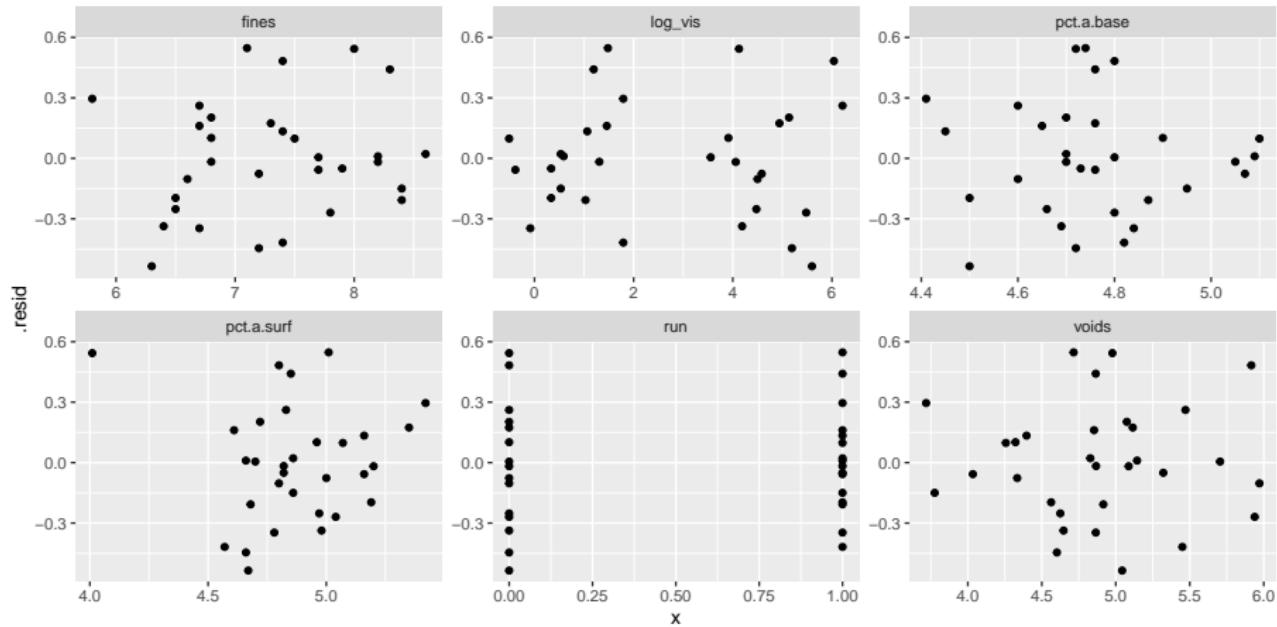
- Do our trick again to put them all on one plot:

```
augment(rut.6)
```

		log(rut.depth)	pct.a.surf	voids	log(viscosity)	.fitted
1		1.91		4.68	4.92	
0.207	0.0772					1.03
2		2.56		5.19	4.56	
0.197	0.111					0.336
3		2.69		4.82	5.32	
0.0503	0.142					0.336
4		2.53		4.85	4.86	
5		2.11		4.86	3.78	
0.150	0.178					0.531
6		2.37		5.16	4.40	
7		1.99		4.82	4.87	
						2.26
						2.09
						2.23
						2.00

Residuals against the x's

g2



Comments

- None of the plots show any sort of pattern. The points all look random on each plot.
- On the plot of fitted values (and on the one of log.viscosity), the points seem to form a “left half” and a “right half” with a gap in the middle. This is not a concern.
- One of the pct.a.surf values is low outlier (4), shows up top left of that plot.
- Only two possible values of run; the points in each group look randomly scattered around 0, with equal spreads.
- Residuals seem to go above zero further than below, suggesting a mild non-normality, but not enough to be a problem.

Variable-selection strategies

- Expert knowledge.
- Backward elimination.
- All possible regressions.
- Taking a variety of models to experts and asking their opinion.
- Use a looser cutoff to eliminate variables in backward elimination (eg. only if P-value greater than 0.20).
- If goal is prediction, eliminating worthless variables less important.
- If goal is understanding, want to eliminate worthless variables where possible.
- Results of variable selection not always reproducible, so caution advised.