

Logistic Regression

Logistic regression

- When response variable is measured/counted, regression can work well.
- But what if response is yes/no, lived/died, success/failure?
- Model *probability* of success.
- Probability must be between 0 and 1; need method that ensures this.
- *Logistic regression* does this. In R, is a *generalized linear model* with binomial “family”:

```
glm(y ~ x, family="binomial")
```

- Begin with simplest case.

Packages

```
library(MASS, exclude = "select")
library(tidyverse)
library(marginaleffects)
library(broom)
library(nnet)
# library(conflicted)
# conflict_prefer("select", "dplyr")
# conflict_prefer("filter", "dplyr")
# conflict_prefer("rename", "dplyr")
# conflict_prefer("summarize", "dplyr")
```

The rats, part 1

- Rats given dose of some poison; either live or die:

dose status

0 lived

1 died

2 lived

3 lived

4 died

5 died

Read in:

```
my_url <- "http://ritsokiguess.site/datafiles/rat.txt"
rats <- read_delim(my_url, " ")
rats
```

```
# A tibble: 6 x 2
```

```
  dose status
  <dbl> <chr>
1     0 lived
2     1 died 
3     2 lived
4     3 lived
5     4 died 
6     5 died
```

This doesn't work

```
status.0 <- glm(status ~ dose, family = "binomial", data = rats)
```

Error in eval(family\$initialize): y values must be 0 <= y <= 1

- Values of response variable (here `status`) must be either:
 - ▶ 1 = “success”, 0 = “failure”
 - ▶ a factor (not text) with two levels.
- The error message doesn't say that the second is a possibility.

Basic logistic regression

- So, make response into a factor first:

```
rats2 <- rats %>% mutate(status = factor(status))  
rats2
```

```
# A tibble: 6 x 2  
  dose status  
  <dbl> <fct>  
1     0 lived  
2     1 died  
3     2 lived  
4     3 lived  
5     4 died  
6     5 died
```

- then fit model:

```
status.1 <- glm(status ~ dose, family = "binomial",  
                  data = rats2)
```

Output

```
summary(status.1)
```

Call:

```
glm(formula = status ~ dose, family = "binomial", data = rats2)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	1.6841	1.7979	0.937	0.349
dose	-0.6736	0.6140	-1.097	0.273

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 8.3178 on 5 degrees of freedom

Residual deviance: 6.7728 on 4 degrees of freedom

AIC: 10.773

Interpreting the output

- Like (multiple) regression, get tests of significance of individual x 's
- Here not significant (only 6 observations).
- “Slope” for dose is negative, meaning that as dose increases, probability of event modelled (survival) decreases.

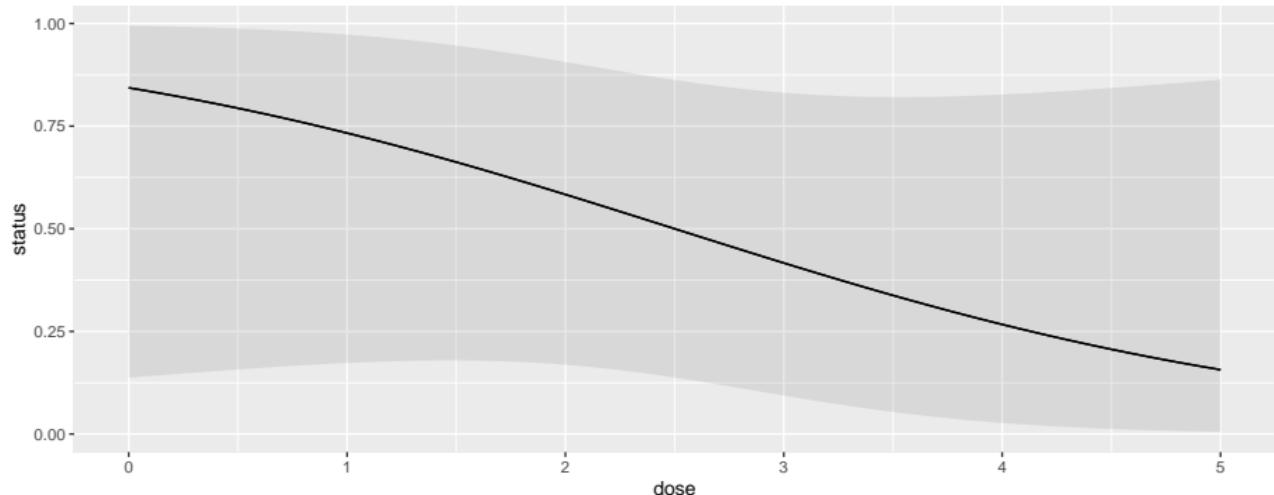
Output part 2: predicted survival probs

```
cbind(predictions(status.1)) %>%
  select(dose, estimate, conf.low, conf.high)
```

	dose	estimate	conf.low	conf.high
1	0	0.8434490	0.137095792	0.9945564
2	1	0.7331122	0.173186479	0.9729896
3	2	0.5834187	0.168847561	0.9061463
4	3	0.4165813	0.093853682	0.8311524
5	4	0.2668878	0.027010413	0.8268135
6	5	0.1565510	0.005443589	0.8629042

On a graph

```
plot_predictions(status.1, condition = "dose")
```



The rats, more

- More realistic: more rats at each dose (say 10).
- Listing each rat on one line makes a big data file.
- Use format below: dose, number of survivals, number of deaths.

dose lived died

0	10	0
1	7	3
2	6	4
3	4	6
4	2	8
5	1	9

- 6 lines of data correspond to 60 actual rats.
- Saved in `rat2.txt`.

These data

```
my_url <- "http://ritsokiguess.site/datafiles/rat2.txt"  
rat2 <- read_delim(my_url, " ")  
rat2
```

```
# A tibble: 6 x 3  
  dose lived died  
  <dbl> <dbl> <dbl>  
1     0     10    0  
2     1      7    3  
3     2      6    4  
4     3      4    6  
5     4      2    8  
6     5      1    9
```

Response matrix:

- Each row contains *multiple* observations.
- Create *two-column* response with cbind:
 - ▶ #survivals in first column,
 - ▶ #deaths in second.

```
with(rat2, cbind(lived, died))
```

	lived	died
[1,]	10	0
[2,]	7	3
[3,]	6	4
[4,]	4	6
[5,]	2	8
[6,]	1	9

Fit logistic regression

- constructing the response in the `glm`:

```
rat2.1 <- glm(cbind(lived, died) ~ dose,  
                family = "binomial", data = rat2)
```

Output

Significant effect of dose now:

```
summary(rat2.1)
```

Call:

```
glm(formula = cbind(lived, died) ~ dose, family = "binomial",
     data = rat2)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)		
(Intercept)	2.3619	0.6719	3.515	0.000439 ***		
dose	-0.9448	0.2351	-4.018	5.87e-05 ***		

Signif. codes:	0 '***'	0.001 '**'	0.01 '*'	0.05 '.'	0.1 ' '	1

(Dispersion parameter for binomial family taken to be 1)

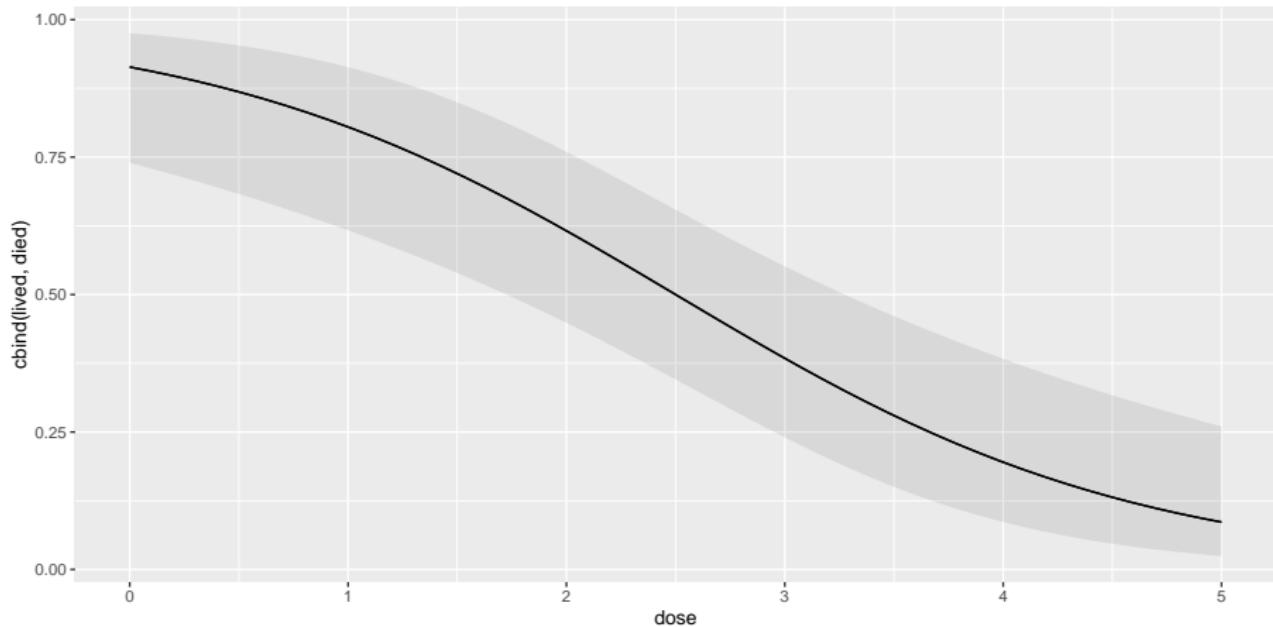
Predicted survival probs

```
new <- datagrid(model = rat2.1, dose = 0:5)
cbind(predictions(rat2.1, newdata = new)) %>%
  select(estimate, dose, conf.low, conf.high)
```

	estimate	dose	conf.low	conf.high
1	0.9138762	0	0.73983042	0.9753671
2	0.8048905	1	0.61695841	0.9135390
3	0.6159474	2	0.44876099	0.7595916
4	0.3840526	3	0.24040837	0.5512390
5	0.1951095	4	0.08646093	0.3830417
6	0.0861238	5	0.02463288	0.2601697

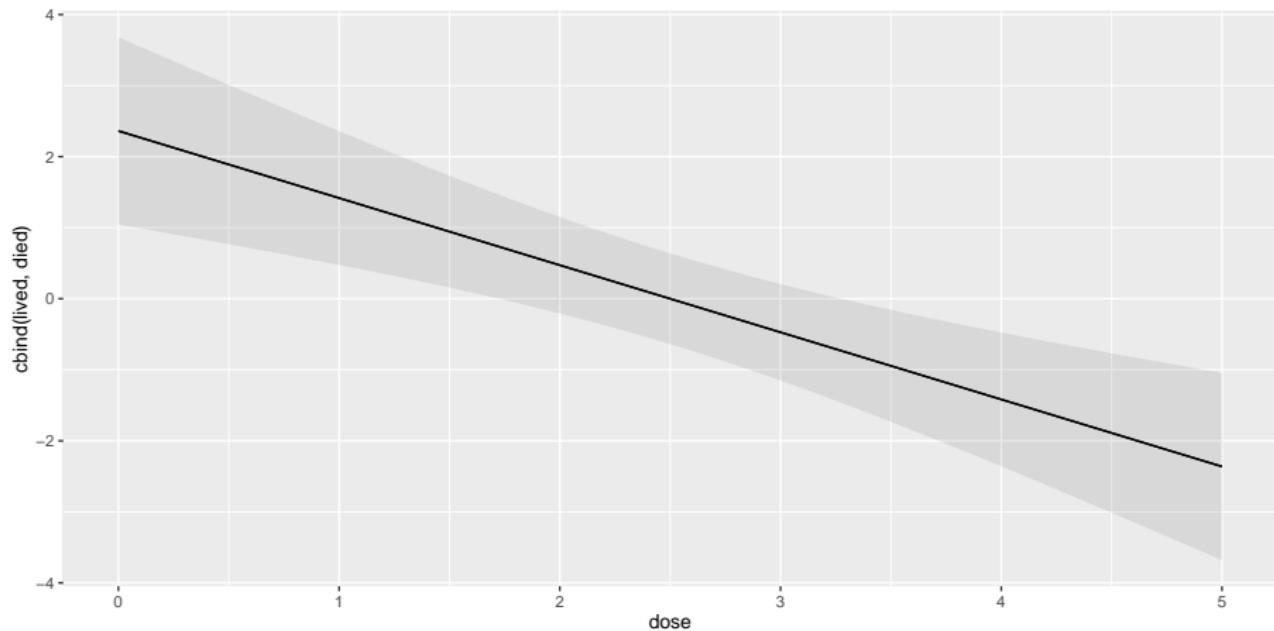
On a picture

```
plot_predictions(rat2.1, condition = "dose")
```



Dose and predicted log-odds

```
plot_predictions(rat2.1, condition = "dose", type = "link")
```



Comments

- Significant effect of dose.
- Effect of larger dose is to *decrease* survival probability (“slope” negative; also see in decreasing predictions.)
- Confidence intervals around prediction narrower (more data).

Multiple logistic regression

- With more than one x , works much like multiple regression.
- Example: study of patients with blood poisoning severe enough to warrant surgery. Relate survival to other potential risk factors.
- Variables, 1=present, 0=absent:
 - ▶ survival (death from sepsis=1), response
 - ▶ shock
 - ▶ malnutrition
 - ▶ alcoholism
 - ▶ age (as numerical variable)
 - ▶ bowel infarction
- See what relates to death.

Read in data

```
my_url <-  
  "http://ritsokiguess.site/datafiles/sepsis.txt"  
sepsis <- read_delim(my_url, " ")  
sepsis
```

```
# A tibble: 106 x 6  
  death shock malnut alcohol age bowelinf  
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
1     0     0     0     0    56     0  
2     0     0     0     0    80     0  
3     0     0     0     0    61     0  
4     0     0     0     0    26     0  
5     0     0     0     0    53     0  
6     1     0     1     0    87     0  
7     0     0     0     0    21     0  
8     1     0     0     1    69     0  
9     0     0     0     0    57     0  
10    0     0     1     0    76     0  
# i 96 more rows
```

Make sure categoricals really are

```
sepsis %>%
  mutate(across(-age, \(x) factor(x))) -> sepsis
```

The data (some)

sepsis

```
# A tibble: 106 x 6
  death shock malnut alcohol    age bowelinf
  <fct> <fct> <fct>   <fct>    <dbl> <fct>
1 0     0     0     0      56  0
2 0     0     0     0      80  0
3 0     0     0     0      61  0
4 0     0     0     0      26  0
5 0     0     0     0      53  0
6 1     0     1     0      87  0
7 0     0     0     0      21  0
8 1     0     0     1      69  0
9 0     0     0     0      57  0
10 0    0     1     0      76  0
# i 96 more rows
```

Fit model

```
sepsis.1 <- glm(death ~ shock + malnut + alcohol + age +  
  bowelinf, family = "binomial", data = sepsis  
)
```

Output part 1

```
summary(sepsis.1)
```

Call:

```
glm(formula = death ~ shock + malnut + alcohol + age + bowelinf,  
     family = "binomial", data = sepsis)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-9.75391	2.54170	-3.838	0.000124	***
shock1	3.67387	1.16481	3.154	0.001610	**
malnut1	1.21658	0.72822	1.671	0.094798	.
alcohol1	3.35488	0.98210	3.416	0.000635	***
age	0.09215	0.03032	3.039	0.002374	**
bowelinf1	2.79759	1.16397	2.403	0.016240	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Or, with tidy (from broom)

```
tidy(sepsis.1)
```

```
# A tibble: 6 x 5
  term      estimate std.error statistic p.value
  <chr>      <dbl>     <dbl>      <dbl>    <dbl>
1 (Intercept) -9.75      2.54     -3.84 0.000124
2 shock1        3.67      1.16      3.15 0.00161 
3 malnut1       1.22      0.728     1.67 0.0948 
4 alcohol1      3.35      0.982     3.42 0.000635
5 age           0.0922    0.0303     3.04 0.00237 
6 bowelinf1     2.80      1.16      2.40 0.0162
```

- All P-values fairly small
- but malnut not significant: remove.

Removing malnut

```
sepsis.2 <- update(sepsis.1, . ~ . - malnut)
summary(sepsis.2)
```

Call:
glm(formula = death ~ shock + alcohol + age + bowelinf, family = "b"
 data = sepsis)

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-8.89459	2.31689	-3.839	0.000124 ***
shock1	3.70119	1.10353	3.354	0.000797 ***
alcohol1	3.18590	0.91725	3.473	0.000514 ***
age	0.08983	0.02922	3.075	0.002106 **
bowelinf1	2.38647	1.07227	2.226	0.026039 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Comments

- Everything significant now.
- Most of the original x 's helped predict death. Only `malnut` seemed not to add anything.
- Removed `malnut` and tried again.
- Everything remaining is significant (though `bowelinf` actually became *less* significant).
- All coefficients are *positive*, so having any of the risk factors (or being older) *increases* risk of death.

Predictions from model without “malnut” 1/2

- A few (rows of original dataframe) chosen “at random”:

```
sepsis %>% slice(c(4, 1, 2, 11, 32)) -> new  
new
```

```
# A tibble: 5 x 6  
  death shock malnut alcohol    age bowelinf  
  <fct> <fct> <fct>   <fct>    <dbl> <fct>  
1 0      0      0       0        26     0  
2 0      0      0       0        56     0  
3 0      0      0       0        80     0  
4 1      0      0       1        66     1  
5 1      0      0       1        49     0
```

Predictions from model without “malnut” 2/2

```
cbind(predictions(sepsis.2, newdata = new)) %>%
  select(estimate, conf.low, conf.high, shock:bowelinf)
```

	estimate	conf.low	conf.high	shock	malnut	alcohol	age	bowelinf
1	0.001415347	6.272642e-05	0.03103047	0	0	0	26	0
2	0.020552383	4.102504e-03	0.09656596	0	0	0	56	0
3	0.153416834	5.606838e-02	0.35603441	0	0	0	80	0
4	0.931290137	5.490986e-01	0.99341482	0	0	1	66	1
5	0.213000997	7.639063e-02	0.46967947	0	0	1	49	0

Comments

- Survival chances pretty good if no risk factors, though decreasing with age.
- Having more than one risk factor reduces survival chances dramatically.
- Usually good job of predicting survival; sometimes death predicted to survive.

Another way to assess effects

of age:

```
new <- datagrid(model = sepsis.2, age = seq(30, 70, 10))  
new
```

	shock	alcohol	bowelinf	age	rowid
1	0	0	0	30	1
2	0	0	0	40	2
3	0	0	0	50	3
4	0	0	0	60	4
5	0	0	0	70	5

Assessing age effect

```
cbind(predictions(sepsis.2, newdata = new)) %>%  
  select(estimate, shock:age)
```

	estimate	shock	alcohol	bowelinf	age
1	0.002026053	0	0	0	30
2	0.004960283	0	0	0	40
3	0.012092515	0	0	0	50
4	0.029179226	0	0	0	60
5	0.068729752	0	0	0	70

Assessing shock effect

```
new <- datagrid(shock = c(0, 1), model = sepsis.2)
new
```

	alcohol	age	bowelinf	shock	rowid
1	0	51.28302	0	0	1
2	0	51.28302	0	1	2

```
cbind(predictions(sepsis.2, newdata = new)) %>%
  select(estimate, alcohol:shock)
```

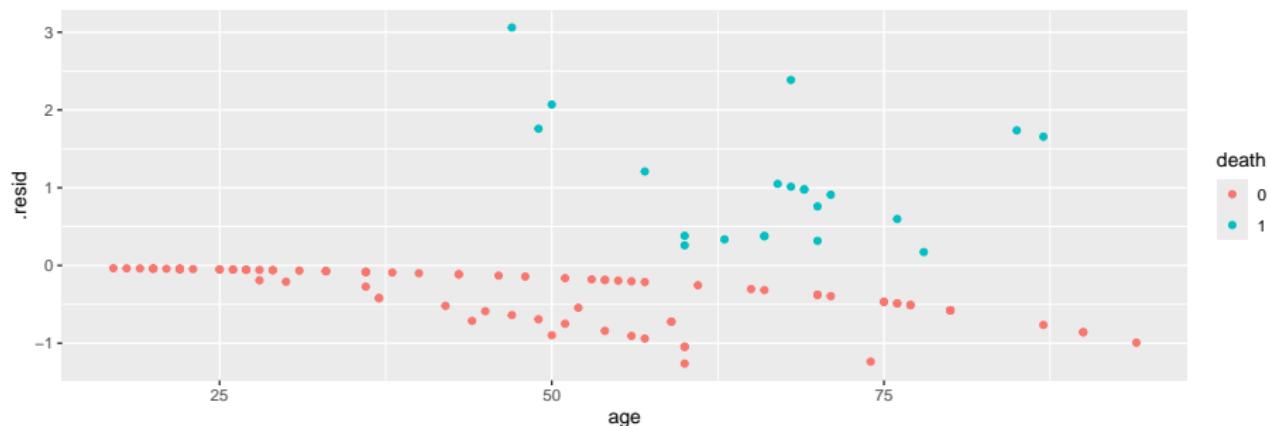
	estimate	alcohol	age	bowelinf	shock
1	0.01354973	0	51.28302	0	0
2	0.35742607	0	51.28302	0	1

Assessing proportionality of odds for age

- An assumption we made is that log-odds of survival depends linearly on age.
- Hard to get your head around, but basic idea is that survival chances go continuously up (or down) with age, instead of (for example) going up and then down.
- In this case, seems reasonable, but should check:

Residuals vs. age

```
sepsis.2 %>% augment(sepsis) %>%
  ggplot(aes(x = age, y = .resid, colour = death)) +
  geom_point()
```



Comments

- No apparent problems overall.
- Confusing “line” across: no risk factors, survived.

Probability and odds

For probability p , odds is $p/(1 - p)$:

Prob	Odds	Log-odds	Words
0.5	$0.5 / 0.5 = 1.00$	0.00	even money
0.1	$0.1 / 0.9 = 0.11$	-2.20	9 to 1
0.4	$0.4 / 0.6 = 0.67$	-0.41	1.5 to 1
0.8	$0.8 / 0.2 = 4.00$	1.39	4 to 1 on

- Gamblers use odds: if you win at 9 to 1 odds, get original stake back plus 9 times the stake.
- Probability has to be between 0 and 1
- Odds between 0 and infinity
- Log-odds can be anything: any log-odds corresponds to valid probability.
- Thus, predict *log-odds of probability* from explanatory variable(s), rather than probability itself.

Odds ratio

- Suppose 90 of 100 men drank wine last week, but only 20 of 100 women.
- Prob of man drinking wine $90/100 = 0.9$, woman $20/100 = 0.2$.
- Odds of man drinking wine $0.9/0.1 = 9$, woman $0.2/0.8 = 0.25$.
- Ratio of odds is $9/0.25 = 36$.
- Way of quantifying difference between men and women: “odds of drinking wine 36 times larger for males than females”.

Sepsis data again

- Recall prediction of probability of death from risk factors:

```
sepsis
```

```
# A tibble: 106 x 6
  death shock malnut alcohol    age bowelinf
  <fct> <fct> <fct>   <fct>    <dbl> <fct>
  1 0     0     0     0      56   0
  2 0     0     0     0      80   0
  3 0     0     0     0      61   0
  4 0     0     0     0      26   0
  5 0     0     0     0      53   0
  6 1     0     1     0      87   0
  7 0     0     0     0      21   0
  8 1     0     0     1      69   0
  9 0     0     0     0      57   0
 10 0    0     1     0      76   0
# i 96 more rows
```

Multiplying the odds

- Can interpret slopes by taking “exp” of them. We ignore intercept.

```
sepsis.2.tidy %>%
  mutate(exp_coeff=exp(estimate)) %>%
  select(term, exp_coeff)
```

```
# A tibble: 5 x 2
  term      exp_coeff
  <chr>     <dbl>
1 (Intercept) 0.000137
2 shock1      40.5
3 alcohol1    24.2
4 age         1.09
5 bowelinf1   10.9
```

Interpretation

```
# A tibble: 5 x 2
  term      exp_coeff
  <chr>        <dbl>
1 (Intercept) 0.000137
2 shock1       40.5
3 alcohol1     24.2
4 age          1.09
5 bowelinf1    10.9
```

- These say “how much do you *multiply* odds of death by for increase of 1 in corresponding risk factor?” Or, what is odds ratio for that factor being 1 (present) vs. 0 (absent)?
- Eg. being alcoholic vs. not increases odds of death by 24 times
- One year older multiplies odds by about 1.1 times. Over 40 years, about $1.09^{40} = 31$ times.

Odds ratio and relative risk

- **Relative risk** is ratio of probabilities.
- Above: 90 of 100 men (0.9) drank wine, 20 of 100 women (0.2).
- Relative risk $0.9/0.2=4.5$. (odds ratio was 36).
- When probabilities small, relative risk and odds ratio similar.
- Eg. prob of man having disease 0.02, woman 0.01.
- Relative risk $0.02/0.01 = 2$.

Odds ratio vs. relative risk

- Odds for men and for women:

```
(od1 <- 0.02 / 0.98) # men
```

```
[1] 0.02040816
```

```
(od2 <- 0.01 / 0.99) # women
```

```
[1] 0.01010101
```

- Odds ratio

```
od1 / od2
```

```
[1] 2.020408
```

- Very close to relative risk of 2.

More than 2 response categories

- With 2 response categories, model the probability of one, and prob of other is one minus that. So doesn't matter which category you model.
- With more than 2 categories, have to think more carefully about the categories: are they
 - ordered*: you can put them in a natural order (like low, medium, high)
 - nominal*: ordering the categories doesn't make sense (like red, green, blue).
- R handles both kinds of response; learn how.

Ordinal response: the miners

- Model probability of being in given category *or lower*.
- Example: coal-miners often suffer disease pneumoconiosis. Likelihood of disease believed to be greater among miners who have worked longer.
- Severity of disease measured on categorical scale: none, moderate, severe.

Miners data

- Data are frequencies:

Exposure	None	Moderate	Severe
5.8	98	0	0
15.0	51	2	1
21.5	34	6	3
27.5	35	5	8
33.5	32	10	9
39.5	23	7	8
46.0	12	6	10
51.5	4	2	5

Reading the data

Data in aligned columns with more than one space between, so:

```
my_url <- "http://ritsokiguess.site/datafiles/miners-tab.txt"  
freqs <- read_table(my_url)
```

The data

```
freqs
```

	Exposure	# A tibble: 8 x 4		
		None	Moderate	Severe
	<dbl>	<dbl>	<dbl>	<dbl>
1	5.8	98	0	0
2	15	51	2	1
3	21.5	34	6	3
4	27.5	35	5	8
5	33.5	32	10	9
6	39.5	23	7	8
7	46	12	6	10
8	51.5	4	2	5

Tidying

```
freqs %>%  
  pivot_longer(-Exposure, names_to = "Severity", values_to = "  
  mutate(Severity = fct_inorder(Severity)) -> miners
```

Result

miners

```
# A tibble: 24 x 3
  Exposure Severity Freq
  <dbl> <fct>    <dbl>
1      5.8 None      98
2      5.8 Moderate     0
3      5.8 Severe       0
4      15  None       51
5      15  Moderate     2
6      15  Severe       1
7     21.5 None       34
8     21.5 Moderate     6
9     21.5 Severe       3
10    27.5 None       35
# i 14 more rows
```

levels(miners\$Severity)

Plot proportions against exposure 1/2

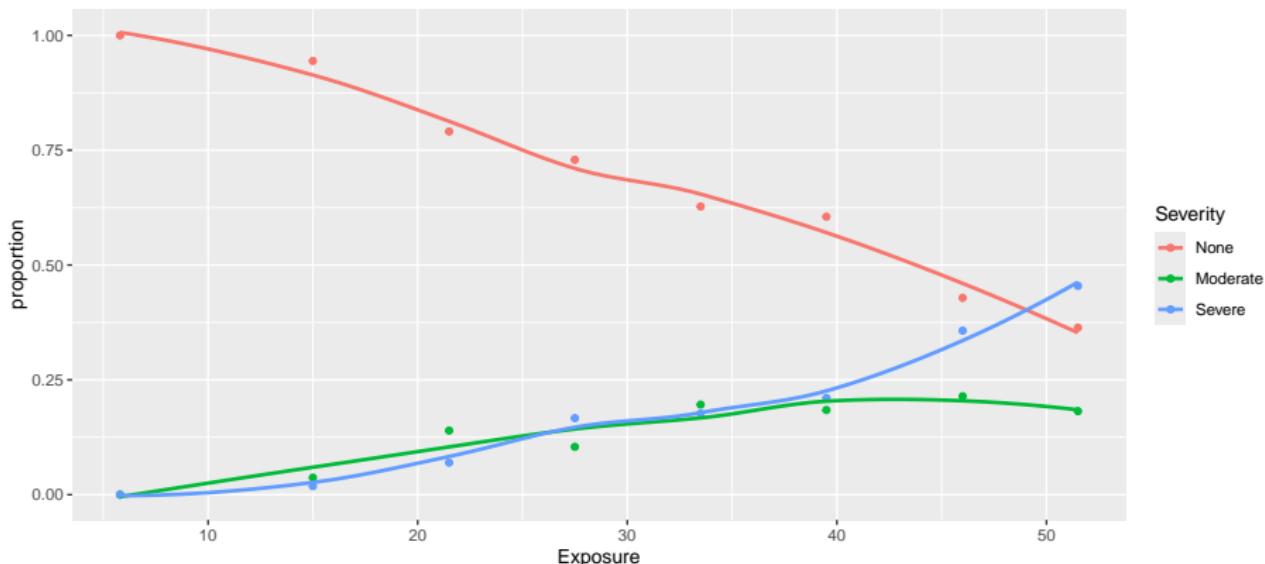
```
miners %>%
  group_by(Exposure) %>%
  mutate(proportion = Freq / sum(Freq)) -> prop
prop
```

```
# A tibble: 24 x 4
# Groups:   Exposure [8]
  Exposure Severity Freq proportion
  <dbl>    <fct>   <dbl>     <dbl>
1      5.8 None      98       1
2      5.8 Moderate  0        0
3      5.8 Severe    0        0
4      15  None      51       0.944
5      15  Moderate  2        0.0370
6      15  Severe    1        0.0185
7     21.5 None      34       0.791
8     21.5 Moderate  6        0.140
9     21.5 Severe    3        0.0698
10     27.5 None      35       0.729
```

i 14 more rows

Plot proportions against exposure 2/2

```
ggplot(prop, aes(x = Exposure, y = proportion,  
                  colour = Severity)) +  
  geom_point() + geom_smooth(se = F)
```



Reminder of data setup

```
miners
```

```
# A tibble: 24 x 3
  Exposure Severity Freq
  <dbl>   <fct>    <dbl>
1      5.8 None        98
2      5.8 Moderate     0
3      5.8 Severe       0
4      15  None         51
5      15  Moderate      2
6      15  Severe        1
7     21.5 None         34
8     21.5 Moderate      6
9     21.5 Severe        3
10    27.5 None         35
# i 14 more rows
```

Fitting ordered logistic model

Use function `polr` from package MASS. Like `glm`.

```
sev.1 <- polr(Severity ~ Exposure,  
  weights = Freq,  
  data = miners  
)
```

Output: not very illuminating

```
sev.1 <- polr(Severity ~ Exposure,  
  weights = Freq,  
  data = miners,  
)
```

```
summary(sev.1)
```

Call:

```
polr(formula = Severity ~ Exposure, data = miners, weights = Freq)
```

Coefficients:

	Value	Std. Error	t value
--	-------	------------	---------

Exposure	0.0959	0.01194	8.034
----------	--------	---------	-------

Intercepts:

	Value	Std. Error	t value
--	-------	------------	---------

None Moderate	3.9558	0.4097	9.6558
---------------	--------	--------	--------

Moderate Severe	4.8690	0.4411	11.0383
-----------------	--------	--------	---------

Residual Deviance: 416.9188

AIC: 422.9188

Does exposure have an effect?

Fit model without Exposure, and compare using anova. Note 1 for model with just intercept:

```
sev.0 <- polr(Severity ~ 1, weights = Freq, data = miners)
anova(sev.0, sev.1)
```

Likelihood ratio tests of ordinal regression models

Response: Severity

	Model	Resid.	df	Resid.	Dev	Test	Df	LR stat.	Pr(Chi)
1		1	369	505.1621					
2	Exposure		368	416.9188	1 vs 2		1	88.24324	0

Exposure definitely has effect on severity of disease.

Another way

- What (if anything) can we drop from model with exposure?

```
drop1(sev.1, test = "Chisq")
```

Single term deletions

Model:

Severity ~ Exposure

	Df	AIC	LRT	Pr(>Chi)
<none>		422.92		

Exposure	1	509.16	88.243	< 2.2e-16	***
----------	---	--------	--------	-----------	-----

Signif. codes:	0	'***'	0.001	'**'	0.01	'*'	0.05	'.'	0.1	' '	1
----------------	---	-------	-------	------	------	-----	------	-----	-----	-----	---

- Nothing. Exposure definitely has effect.

Predicted probabilities 1/2

```
freqs %>% select(Exposure) -> new  
new
```

```
# A tibble: 8 x 1
```

```
  Exposure
```

```
  <dbl>
```

```
1      5.8
2     15
3    21.5
4   27.5
5   33.5
6   39.5
7   46
8  51.5
```

Predicted probabilities 2/2

```
cbind(predictions(sev.1, newdata = new)) %>%
  select(group, estimate, Exposure) %>%
  pivot_wider(names_from = group, values_from = estimate)
```

A tibble: 8 x 4

	Exposure	None	Moderate	Severe
	<dbl>	<dbl>	<dbl>	<dbl>
1	5.8	0.968	0.0191	0.0132
2	15	0.925	0.0433	0.0314
3	21.5	0.869	0.0739	0.0569
4	27.5	0.789	0.114	0.0969
5	33.5	0.678	0.162	0.160
6	39.5	0.542	0.205	0.253
7	46	0.388	0.224	0.388
8	51.5	0.272	0.210	0.517

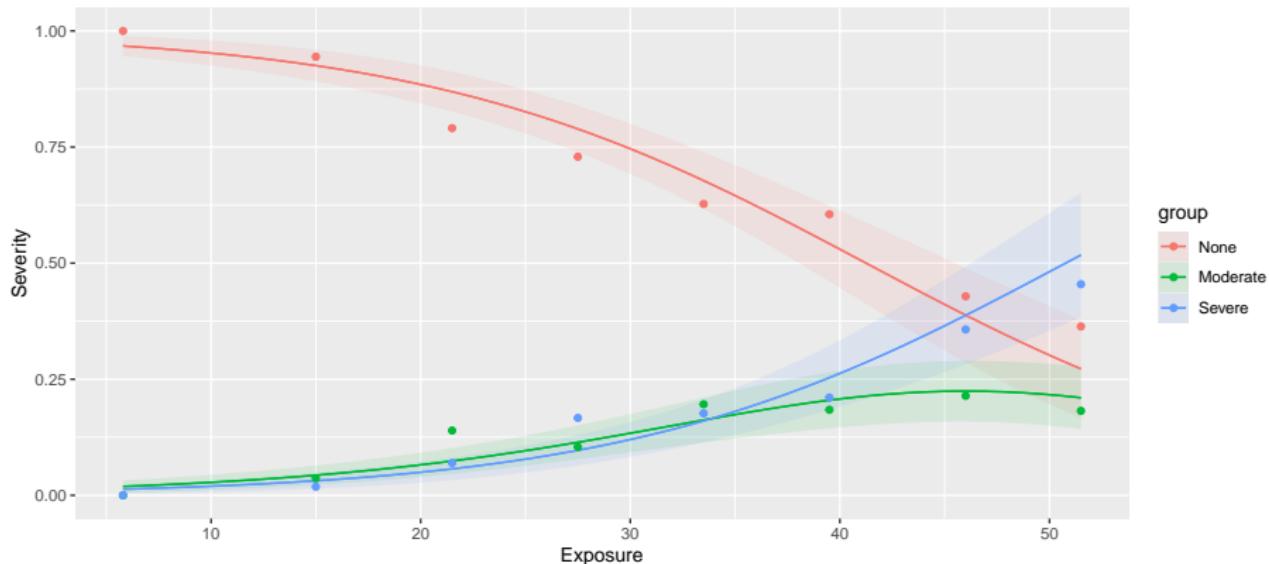
Plot of predicted probabilities

- Wider for looking at, longer for graph:

```
plot_predictions(model = sev.1,
                  condition = c("Exposure", "group"),
                  type = "probs") +
  geom_point(data = prop, aes(x = Exposure, y = proportion,
                               colour = Severity)) -> ggg
```

The graph

ggg



Comments

- Model appears to match data well enough.
- As exposure goes up, prob of None goes down, Severe goes up (sharply for high exposure).
- So more exposure means worse disease.

Unordered responses

- With unordered (nominal) responses, can use *generalized logit*.
- Example: 735 people, record age and sex (male 0, female 1), which of 3 brands of some product preferred.
- Data in `mlogit.csv` separated by commas (so `read_csv` will work):

```
my_url <- "http://ritsokiguess.site/datafiles/mlogit.csv"  
brandpref <- read_csv(my_url)
```

The data (some)

```
brandpref
```

```
# A tibble: 735 x 3
  brand   sex   age
  <dbl> <dbl> <dbl>
1     1     0    24
2     1     0    26
3     1     0    26
4     1     1    27
5     1     1    27
6     3     1    27
7     1     0    27
8     1     0    27
9     1     1    27
10    1     0    27
# i 725 more rows
```

Bashing into shape

- sex and brand not meaningful as numbers, so turn into factors:

```
brandpref %>%
  mutate(sex = ifelse(sex == 1, "female", "male"),
        sex = factor(sex),
        brand = factor(brand)
  ) -> brandpref
```

```
brandpref
```

```
# A tibble: 735 x 3
  brand sex     age
  <fct> <fct>   <dbl>
1 1     male    24
2 1     male    26
3 1     male    26
4 1     female  27
5 1     female  27
6 3     female  27
7 1     male    27
```

Fitting model

- We use `multinom` from package `nnet`. Works like `polr`.

```
library(nnet)
# levels(brandpref$sex)

brands.1 <- multinom(brand ~ age + sex, data = brandpref)

# weights: 12 (6 variable)
initial value 807.480032
iter 10 value 702.990572
final value 702.970704
converged
```

- summary output not helpful.

Can we drop anything?

- Unfortunately drop1 seems not to work:

```
drop1(brands.1, test = "Chisq", trace = 0)
```

trying - age

Error in if (trace) {: argument is not interpretable as logical

- So, fall back on fitting model without what you want to test, and comparing using anova.

Do age/sex help predict brand? 1/3

Fit models without each of age and sex:

```
brands.2 <- multinom(brand ~ age, data = brandpref)
```

```
# weights: 9 (4 variable)
initial value 807.480032
iter 10 value 706.796323
iter 10 value 706.796322
final value 706.796322
converged
```

```
brands.3 <- multinom(brand ~ sex, data = brandpref)
```

```
# weights: 9 (4 variable)
initial value 807.480032
final value 791.861266
converged
```

Do age/sex help predict brand? 2/3

```
anova(brands.2, brands.1)
```

Likelihood ratio tests of Multinomial Models

Response: brand

	Model	Resid.	df	Resid.	Dev	Test	Df	LR stat.	Pr(Chi)
1	age		1466		1413.593				
2	age + sex		1464		1405.941	1 vs 2	2	7.651236	0.02180496

Do age/sex help predict brand? 3/3

```
anova(brands.3, brands.1)
```

Likelihood ratio tests of Multinomial Models

Response: brand

	Model	Resid.	df	Resid.	Dev	Test	Df	LR stat.	Pr(Chi)
1	sex		1466		1583.723				
2	age + sex		1464		1405.941	1 vs 2	2	177.7811	0

Comments

- age definitely significant (second anova)
- sex significant also (first anova), though P-value less dramatic
- Keep both.
- Expect to see a large effect of age, and a smaller one of sex.

Another way to build model

- Start from model with everything and run step:

```
step(brands.1, trace = 0)
```

trying - age

trying - sex

Call:

```
multinom(formula = brand ~ age + sex)
```

Coefficients:

	(Intercept)	age	sexmale
2	-11.25127	0.3682202	-0.5237736
3	-22.25571	0.6859149	-0.4658215

Residual Deviance: 1405.941

AIC: 1417.941

- Final model contains both age and sex so neither could be removed.

Making predictions

Find age 5-number summary, and the two sexes:

```
summary(brandpref)
```

brand	sex	age
1:207	female:466	Min. :24.0
2:307	male :269	1st Qu.:32.0
3:221		Median :32.0
		Mean :32.9
		3rd Qu.:34.0
		Max. :38.0

Space the ages out a bit for prediction (see over).

Combinations

```
new <- datagrid(age = seq(24, 40, 4), # cover age range  
                  sex = c("female", "male"), model = brands.1)  
new
```

	age	sex	rowid
1	24	female	1
2	24	male	2
3	28	female	3
4	28	male	4
5	32	female	5
6	32	male	6
7	36	female	7
8	36	male	8
9	40	female	9
10	40	male	10

The predictions

```
cbind(predictions(brands.1, newdata = new)) %>%
  select(group, estimate, age, sex) %>%
  pivot_wider(names_from = group, values_from = estimate)
```

A tibble: 10 x 5

	age	sex	`1`	`2`	`3`
	<dbl>	<fct>	<dbl>	<dbl>	<dbl>
1	24	female	0.915	0.0819	0.00279
2	24	male	0.948	0.0502	0.00181
3	28	female	0.696	0.271	0.0329
4	28	male	0.793	0.183	0.0236
5	32	female	0.291	0.495	0.214
6	32	male	0.405	0.408	0.187
7	36	female	0.0503	0.374	0.576
8	36	male	0.0795	0.350	0.571
9	40	female	0.00473	0.153	0.842
10	40	male	0.00759	0.146	0.847

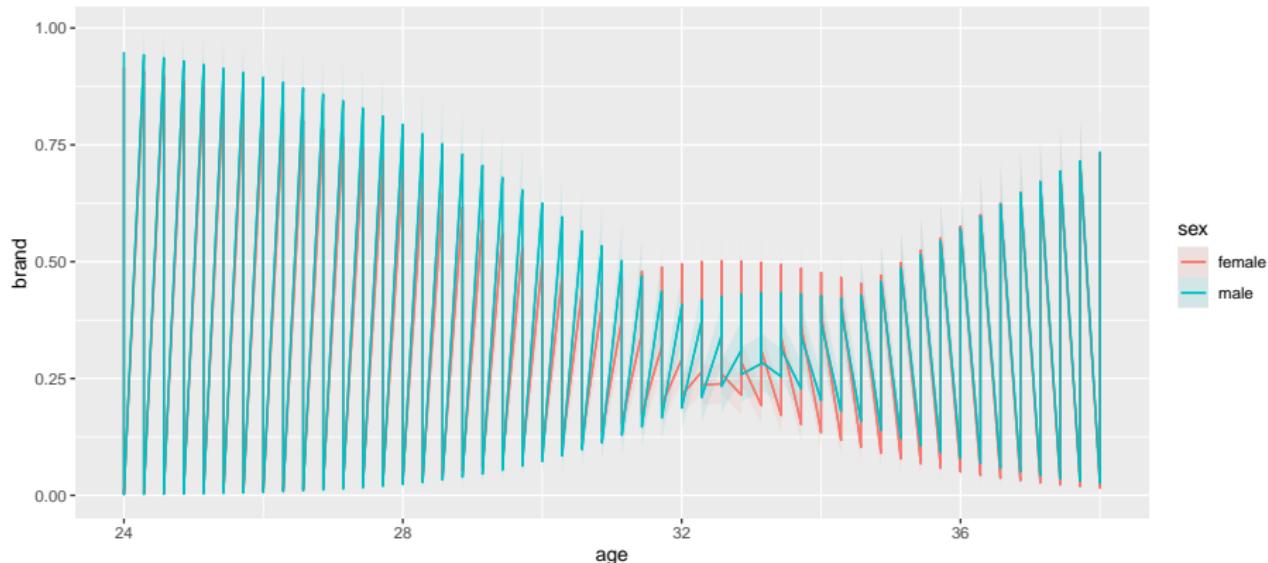
Comments

- Young males prefer brand 1, but older males prefer brand 3.
- Females similar, but like brand 1 less and brand 2 more.
- A clear brand effect, but the sex effect is less clear.

Making a plot

- I thought `plot_predictions` doesn't work as we want, but I was (sort of) wrong about that:

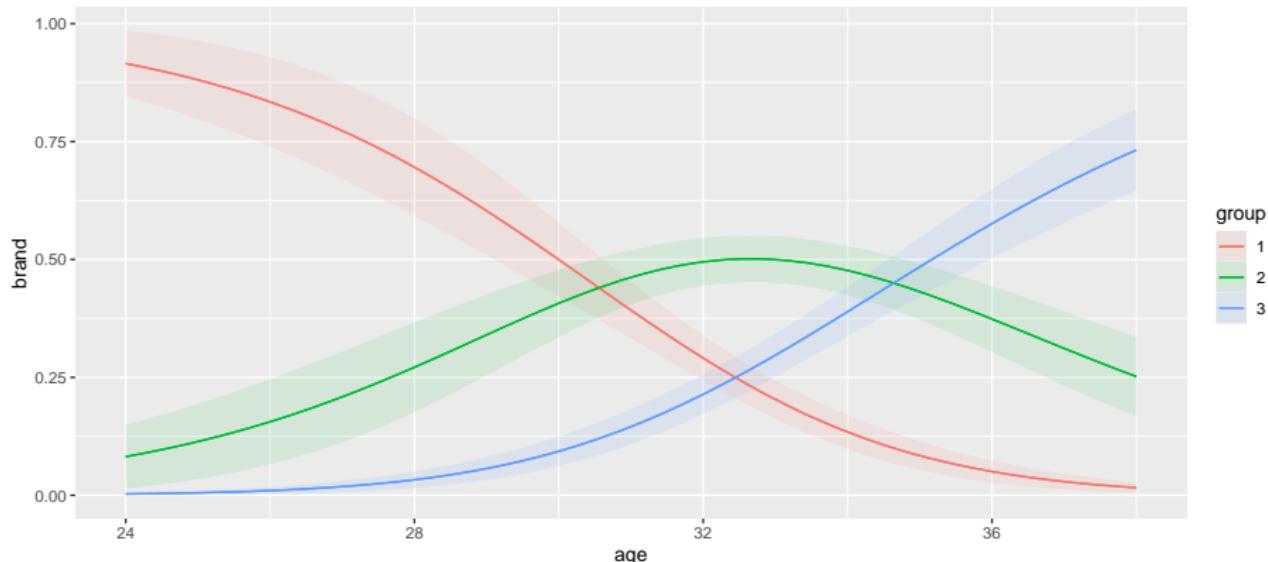
```
plot_predictions(brands.1, condition = c("age", "sex"),
                 type = "probs")
```



Making it go

- We have to include group in the condition:

```
plot_predictions(brands.1, condition = c("age", "group"))
```



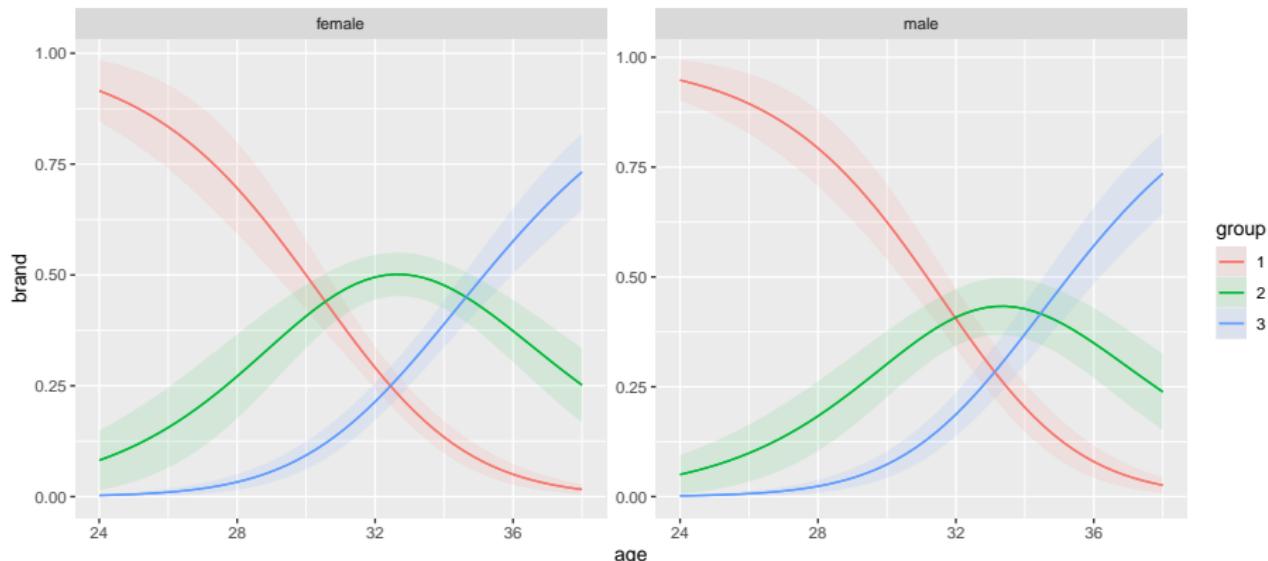
Comments

- This picks the most common sex in the data (females).
- See younger females prefer brand 1, older ones preferring brand 3.

For each sex

If we add the other variable to the *end*, we get facets for sex:

```
plot_predictions(brands.1, condition = c("age", "group", "sex"))
```



Not actually much difference between males and females.

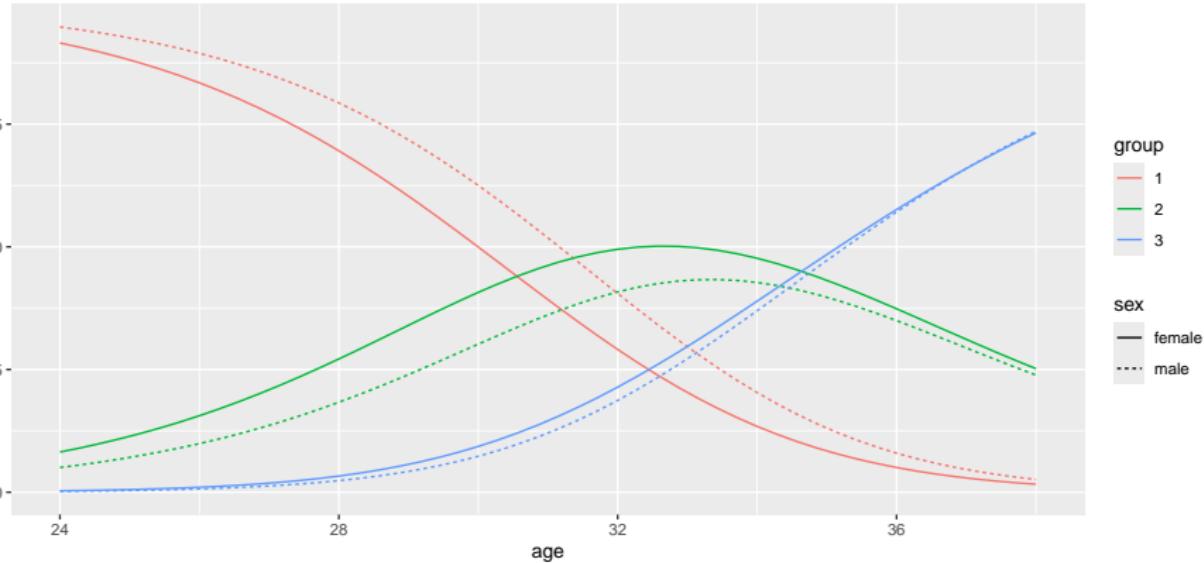
A better graph

- but the male-female difference *was* significant. How?
- *don't* actually plot the graph, then plot the right things:

```
plot_predictions(brands.1,
                 condition = c("age", "brand", "sex"),
                 type = "probs", draw = FALSE) %>%
ggplot(aes(x = age, y = estimate, colour = group,
           linetype = sex)) +
geom_line() -> g
```

The graph

gg



Digesting the plot

- Brand vs. age: younger people (of both genders) prefer brand 1, but older people (of both genders) prefer brand 3. (Explains significant age effect.)
- Brand vs. sex: females (solid) like brand 1 less than males (dashed), like brand 2 more (for all ages).
- Not much brand difference between genders (solid and dashed lines of same colours close), but enough to be significant.
- Model didn't include interaction, so modelled effect of gender on brand same for each age, modelled effect of age same for each gender. (See also later.)

Alternative data format

Summarize all people of same brand preference, same sex, same age on one line of data file with frequency on end:

1 0 24 1

1 0 26 2

1 0 27 4

1 0 28 4

1 0 29 7

1 0 30 3

...

Whole data set in 65 lines not 735! But how?

Getting alternative data format

```
brandpref %>%
  group_by(age, sex, brand) %>%
  summarize(Freq = n()) %>%
  ungroup() -> b
```

b

```
# A tibble: 65 x 4
  age   sex   brand Freq
  <dbl> <fct> <fct> <int>
1 24    male   1      1
2 26    male   1      2
3 27    female 1      4
4 27    female 3      1
5 27    male   1      4
6 28    female 1      6
7 28    female 2      2
8 28    female 3      1
9 28    female 4      1
```

Fitting models, almost the same

- Just have to remember weights to incorporate frequencies.
- Otherwise `multinom` assumes you have just 1 obs on each line!
- Again turn (numerical) sex and brand into factors:

```
b %>%
  mutate(sex = factor(sex)) %>%
  mutate(brand = factor(brand)) -> bf
b.1 <- multinom(brand ~ age + sex, data = bf, weights = Freq)
b.2 <- multinom(brand ~ age, data = bf, weights = Freq)
```

P-value for sex identical

```
anova(b.2, b.1)
```

Likelihood ratio tests of Multinomial Models

Response: brand

	Model	Resid.	df	Resid.	Dev	Test	Df	LR stat.	Pr(Chi)
1	age		126		1413.593				
2	age + sex		124		1405.941	1 vs 2	2	7.651236	0.02180496

Same P-value as before, so we haven't changed anything important.

Trying interaction between age and sex

```
brands.4 <- update(brands.1, . ~ . + age:sex)
```

```
anova(brands.1, brands.4)
```

Likelihood ratio tests of Multinomial Models

Response: brand

	Model	Resid.	df	Resid.	Dev	Test	Df	LR stat.	Pr(Chi)
1	age + sex		1464		1405.941				
2	age + sex + age:sex		1462		1405.142	1 vs 2	2	0.7996223	0.6704466

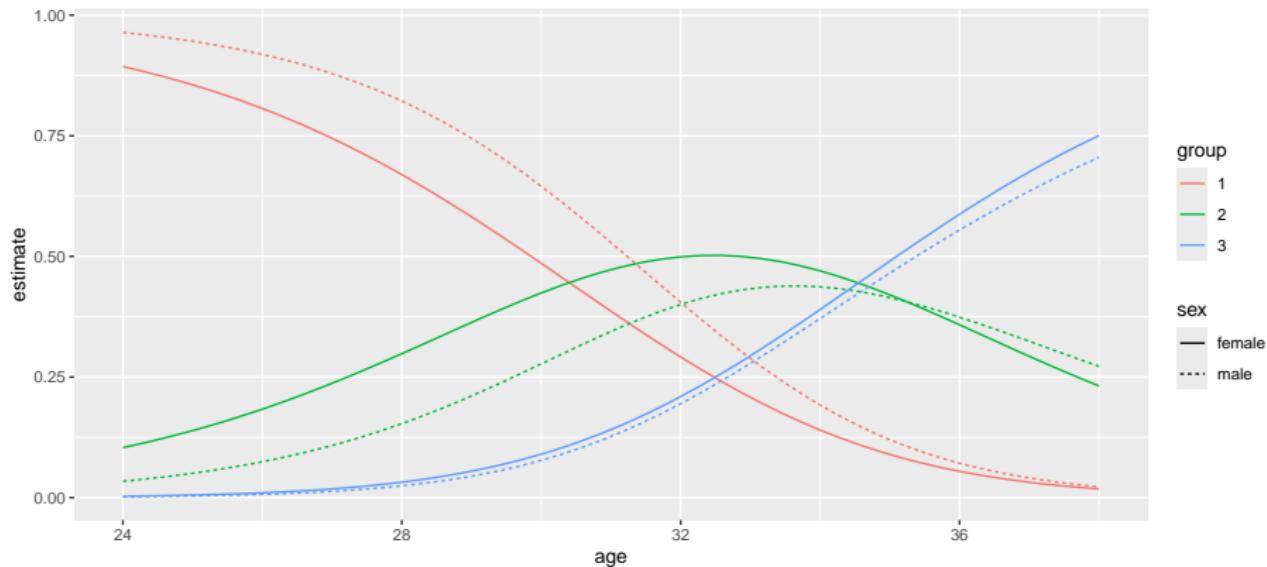
- No evidence that effect of age on brand preference differs for the two genders.

Make graph again

```
plot_predictions(brands.4,
                 condition = c("age", "brand", "sex"),
                 type = "probs", draw = FALSE) %>%
ggplot(aes(x = age, y = estimate, colour = group,
           linetype = sex)) +
geom_line() -> g4
```

Not much difference in the graph

g4



Compare model without interaction

gg

