Repeated measures analysis

Repeated measures

- More than one response *measurement* for each subject, same thing at different times
- Generalization of matched pairs ("matched triples", etc.).
- Expect measurements on same subject to be correlated, so assumptions of independence will fail.
- *Repeated measures. Profile analysis* uses `Manova` (set up).
- Another approach uses *mixed models* (random effects).
- Variation: each subject does all treatments at different times (called *crossover design*).

Packages

```
library(car)
library(tidyverse)
library(lme4) # for mixed models later
```

Example: histamine in dogs

- 8 dogs take part in experiment.
- Dogs randomized to one of 2 different drugs.
- Response: log of blood concentration of histamine 0, 1, 3 and 5 minutes after taking drug. (Repeated measures.)
- Data in dogs.txt, column-aligned.

Read in data

```
my_url <- "http://ritsokiguess.site/datafiles/dogs.txt"
dogs <- read_table(my_url)
dogs
```

```
# A tibble: 8 x 7
  dog   drug          x       lh0   lh1   lh3   lh5
  <chr> <chr>         <chr> <dbl> <dbl> <dbl> <dbl>
1 A     Morphine      N     -3.22 -1.61 -2.3  -2.53
2 B     Morphine      N     -3.91 -2.81 -3.91 -3.91
3 C     Morphine      N     -2.66  0.34 -0.73 -1.43
4 D     Morphine      N     -1.77 -0.56 -1.05 -1.43
5 E     Trimethaphan  N     -3.51 -0.48 -1.17 -1.51
6 F     Trimethaphan  N     -3.51  0.05 -0.31 -0.51
7 G     Trimethaphan  N     -2.66 -0.19  0.07 -0.22
8 H     Trimethaphan  N     -2.41  1.14  0.72  0.21
```

Setting things up

```
response <- with(dogs, cbind(lh0, lh1, lh3, lh5))
response
```

```
        lh0   lh1   lh3   lh5
[1,] -3.22 -1.61 -2.30 -2.53
[2,] -3.91 -2.81 -3.91 -3.91
[3,] -2.66  0.34 -0.73 -1.43
[4,] -1.77 -0.56 -1.05 -1.43
[5,] -3.51 -0.48 -1.17 -1.51
[6,] -3.51  0.05 -0.31 -0.51
[7,] -2.66 -0.19  0.07 -0.22
[8,] -2.41  1.14  0.72  0.21
```

Another way to make response

```
dogs %>% select(starts_with("lh")) %>%
  as.matrix() -> response
response
```

```
        lh0    lh1    lh3    lh5
[1,]  -3.22  -1.61  -2.30  -2.53
[2,]  -3.91  -2.81  -3.91  -3.91
[3,]  -2.66   0.34  -0.73  -1.43
[4,]  -1.77  -0.56  -1.05  -1.43
[5,]  -3.51  -0.48  -1.17  -1.51
[6,]  -3.51   0.05  -0.31  -0.51
[7,]  -2.66  -0.19   0.07  -0.22
[8,]  -2.41   1.14   0.72   0.21
```

The repeated measures MANOVA

Get list of response variable names; we call them `times`. Save in data
frame.

```
times <- colnames(response)
times
```

```
[1] "lh0" "lh1" "lh3" "lh5"
```

```
times.df <- data.frame(times=factor(times))
times.df
```

```
  times
1   lh0
2   lh1
3   lh3
4   lh5
```

Fitting the model

```
dogs.1 <- lm(response ~ drug, data = dogs)
dogs.2 <- Manova(dogs.1,
  idata = times.df,
  idesign = ~times
)
```

The output (there is a lot)

- normally you just run

```
summary(dogs.2)
```

and pull out what you need to answer the question.

- But you can grab just individual pieces as shown below:

```
names(summary(dogs.2))
```

```
[1] "type"               "repeated"          "multivariate.tests"
[4] "univariate.tests"   "pval.adjustments"  "sphericity.tests"
[7] "SSPE"
```

What there is here

- three sets of tests, for
  - times; drug; their interaction
- two *types* of test for each of these:
  - univariate; multivariate
- univariate is more powerful *if* it applies; if it doesn't, can make adjustments to it

Sphericity

- The thing that decides whether the univariate tests apply is called "sphericity".
- This holds if the outcomes have equal variance (to each other) and have the same (positive) correlation across subjects.
- Tested using Mauchly's test (part of output)
- If sphericity rejected, there are adjustments to the univariate P-values due to Huynh-Feldt and Greenhouse-Geisser. Huynh-Feldt better if responses not actually normal (safer).

Sphericity tests

```
summary(dogs.2)$sphericity.tests
```

```
           Test statistic  p-value
times              0.12334 0.084567
drug:times         0.12334 0.084567
```

Sphericity is not rejected; proceed to univariate tests.

Univariate tests

```
summary(dogs.2)$univariate.tests
```

```
            Sum Sq num Df Error SS den Df F value    Pr(>F)
(Intercept) 71.342      1  22.1026      6 19.3664  0.004565 **
drug        11.520      1  22.1026      6  3.1272  0.127406
times       26.160      3   2.2534     18 69.6546 4.215e-10 ***
drug:times   5.111      3   2.2534     18 13.6095 7.050e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Significant interaction between drug and time: the pattern of log-histamine over time is different for the different drugs.

If sphericity had been rejected

then we would use the H-F adjusted P-values:

```
summary(dogs.2)$pval.adjustments
```

```
             GG eps   Pr(>F[GG])     HF eps   Pr(>F[HF])
times      0.5261798 3.744618e-06 0.6822614 1.843418e-07
drug:times 0.5261798 2.348896e-03 0.6822614 7.307096e-04
attr(,"na.action")
(Intercept)          drug
          1             2
attr(,"class")
[1] "omit"
```

In this case (because sphericity was not rejected), these are very similar to the ones from the univariate tests, and the conclusion (significant interaction) was the same.

Comments

- If the interaction had not been significant:
  - ▶ cannot remove interaction with time
  - ▶ so look at univariate (or adjusted for sphericity) tests of main effects in model with non-significant interaction

Next

- investigate interaction with graph
- but dataframe has several observations per line ("wide").
- Plotting works with data in "long format": one response per line.
- The responses are log-histamine at different times, labelled lh-something. Call them all lh and put them in one column, with the time they belong to labelled.

Running `pivot_longer`, try 1

```
dogs %>% pivot_longer(starts_with("lh"),
                      names_to = "time", values_to = "lh")
```

```
# A tibble: 32 x 5
   dog   drug      x     time     lh
   <chr> <chr>    <chr> <chr> <dbl>
 1 A     Morphine N     lh0   -3.22
 2 A     Morphine N     lh1   -1.61
 3 A     Morphine N     lh3   -2.3
 4 A     Morphine N     lh5   -2.53
 5 B     Morphine N     lh0   -3.91
 6 B     Morphine N     lh1   -2.81
 7 B     Morphine N     lh3   -3.91
 8 B     Morphine N     lh5   -3.91
 9 C     Morphine N     lh0   -2.66
10 C     Morphine N     lh1    0.34
# i 22 more rows
```

Getting the times

Not quite right: want new variable containing just number in `time`:
`parse_number`. (Top 5 rows shown.)

```
dogs %>%
  pivot_longer(starts_with("lh"),
               names_to = "timex", values_to = "lh") %>%
  mutate(time = parse_number(timex))
```

```
# A tibble: 5 x 6
  dog   drug     x     timex    lh  time
  <chr> <chr>    <chr> <chr> <dbl> <dbl>
1 A     Morphine N     lh0   -3.22     0
2 A     Morphine N     lh1   -1.61     1
3 A     Morphine N     lh3   -2.3      3
4 A     Morphine N     lh5   -2.53     5
5 B     Morphine N     lh0   -3.91     0
```

What I did differently

- I realized that `pivot_longer` was going to produce something like `lh1`, which I needed to do something further with, so this time I gave it a temporary name `timex` (which we actually *do* use later).
- This enabled me to use the name `time` for the actual numeric time.
- This works now, so next save into a new data frame `dogs.long`.

Saving

```
dogs %>%
  pivot_longer(starts_with("lh"),
               names_to = "timex", values_to = "lh") %>%
  mutate(time = parse_number(timex)) -> dogs.long
```
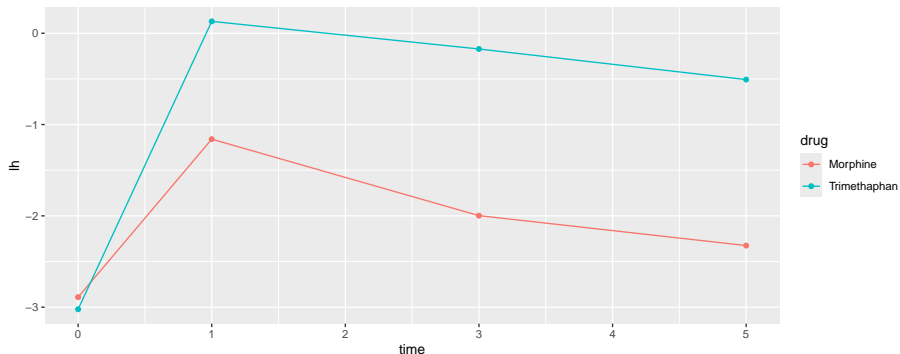
Comments

This says:

- Take data frame dogs, and then:
- Combine the columns `lh0` through `lh5` into one column called `lh`, with the column that each `lh` value originally came from labelled by `timex`, and then:
- Pull out numeric values in `timex`, saving in `time` and then:
- save the result in a data frame `dogs.long`.

Interaction plot

```
ggplot(dogs.long, aes(x = time, y = lh,
                      colour = drug, group = drug)) +
  stat_summary(fun = mean, geom = "point") +
  stat_summary(fun = mean, geom = "line")
```

Comments

- Plot mean `lh` value at each time, joining points on same drug by lines.
- drugs same at time 0
- after that, Trimethaphan higher than Morphine.
- Effect of drug not consistent over time: significant interaction.

Take out time zero

- Lines on interaction plot would then be parallel, and so interaction should no longer be significant.
- Go back to original "wide" dogs data frame.

```
response <- with(dogs, cbind(lh1, lh3, lh5)) # excl time 0
dogs.1 <- lm(response ~ drug, data = dogs)
times <- colnames(response)
times.df <- data.frame(times=factor(times))
dogs.2 <- Manova(dogs.1,
  idata = times.df,
  idesign = ~times
)
```

Results (univariate)

```
summary(dogs.2)$sphericity.tests
```

```
         Test statistic p-value
times          0.57597 0.25176
drug:times     0.57597 0.25176
```

```
# summary(dogs.2)$pval.adjustments
summary(dogs.2)$univariate.tests
```

```
            Sum Sq num Df Error SS den Df F value    Pr(>F)
(Intercept) 24.2607      1  20.1874      6  7.2106   0.03628 *
drug        16.2197      1  20.1874      6  4.8207   0.07053 .
times        3.3250      2   0.7301     12 27.3251 3.406e-05 ***
drug:times   0.3764      2   0.7301     12  3.0929   0.08254 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Comments

- sphericity: no problem (P-value 0.25)
- univariate test for interaction no longer significant (P-value 0.082)
- look at main effects:
  - strong significance of time, even after taking out time 0
 - actually *not* significant drug effect, despite interaction plot

Non-significant drug effect reasonable?

- Plot *actual data*: lh against days, labelling observations by drug: "spaghetti plot".
- Uses long data frame:
  - Plot (time, lh) points coloured by *drug*
  - connecting measurements for each *dog* by lines.
  - Hence, group = dog, but colour = drug:

```
ggplot(dogs.long, aes(x = time, y = lh,
  colour = drug, group = dog)) +
  geom_point() + geom_line() -> g
```

The spaghetti plot

g
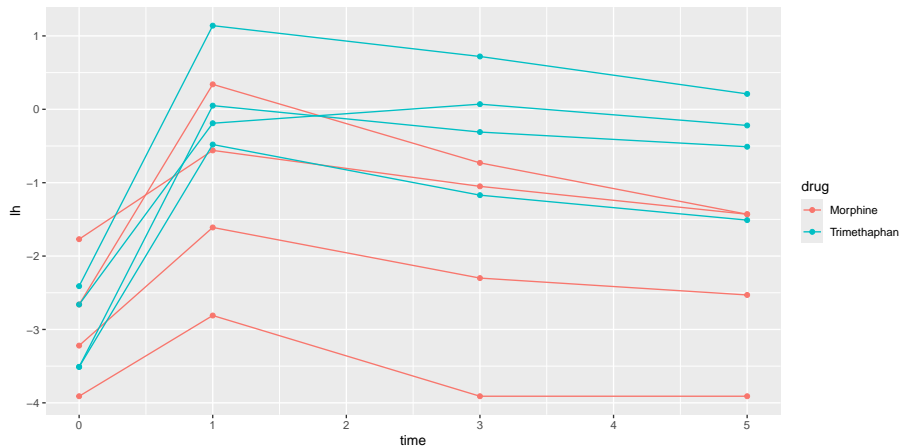
Comments

- For each dog over time, gradual decrease in log-histamine from time 1: significant time effect after we took out time 0.
- Pattern about same for each dog, regardless of drug, hence non-significant interaction.
- Most trimethaphan dogs (blue) have higher log-histamine throughout (time 1 and after), some morphine dogs (red) have lower.
- *But* two morphine dogs have log-histamine profiles like trimethaphan dogs. This ambiguity probably why `drug` effect not quite significant.

Mixed models

- Another way to fit repeated measures
- Subjects (on whom repeated measures taken) are *random sample of all possible subjects* (random effects)
- Times and treatments are *the only ones we care about* (fixed effects)
- Use package `lme4` function `lmer` (like `lm` in some ways)
- Uses long-format "tidy" data

Fitting the model (uses `lme4`)

```
# dogs.long including time zero with categorical timex
dogs.3 <- lmer(lh ~ drug * timex + (1|dog), data=dogs.long)
```

- note specification of random effect: each dog has "random intercept" that moves log-histamine up or down for that dog over all times

What can we drop?

- using drop1:

```
drop1(dogs.3, test="Chisq")
```

```
Single term deletions

Model:
lh ~ drug * timex + (1 | dog)
          npar   AIC     LRT   Pr(Chi)
<none>          62.167
drug:timex   3 84.589  28.422  2.962e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Interaction very significant. Including time zero, the pattern of
  log-histamine over time is different for the two drugs (as we found
  before).

Omitting time zero

Let's pretend we are working at $\alpha = 0.01$:

```
dogs.long %>% filter(timex != "lh0") -> dogs.long.no0
dogs.4 <- lmer(lh ~ drug * timex + (1|dog), data=dogs.long.no0
drop1(dogs.4, test = "Chisq")
```

```
Single term deletions

Model:
lh ~ drug * timex + (1 | dog)
           npar   AIC    LRT Pr(Chi)
<none>            42.119
drug:timex    2 44.771 6.6518 0.03594 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Interaction is not quite significant at $\alpha = 0.01$. So we could remove it.

Removing the interaction

```
dogs.5 <- update(dogs.4, . ~ . - drug:timex)
drop1(dogs.5, test = "Chisq")
```

```
Single term deletions

Model:
lh ~ drug + timex + (1 | dog)
       npar    AIC    LRT   Pr(Chi)
<none>        44.771
drug      1 47.489  4.7176  0.02985 *
timex     2 62.972 22.2011 1.51e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Definitely an effect of time, but drug is not quite significant (at $\alpha = 0.01$).
- More or less same conclusions as from MANOVA.

The exercise data

- 30 people took part in an exercise study.
- Each subject randomly assigned to one of two diets ("low fat" or "non-low fat") and to one of three exercise programs ("at rest", "walking", "running").
- $2 \times 3 = 6$ experimental treatments, and thus each one replicated $30/6 = 5$ times. (Two-way ANOVA, so far?)
- However, each subject had pulse rate measured at three different times (1, 15 and 30 minutes after starting their exercise), so have repeated measures.

Reading the data

Separated by *tabs*:

```
url <- "http://ritsokiguess.site/datafiles/exercise2.txt"
exercise.long <- read_tsv(url)
exercise.long %>% slice(1:7) # top 7 rows
```

```
# A tibble: 7 x 5
     id diet      exertype pulse time
  <dbl> <chr>     <chr>    <dbl> <chr>
1     1 nonlowfat atrest      85 min01
2     1 nonlowfat atrest      85 min15
3     1 nonlowfat atrest      88 min30
4     2 nonlowfat atrest      90 min01
5     2 nonlowfat atrest      92 min15
6     2 nonlowfat atrest      93 min30
7     3 nonlowfat atrest      97 min01
```

Comments

- "Long format", usually what we want.
- But for repeated measures analysis, we want *wide* format!
- Keep track of which is which:
  - ▶ `Manova` analysis: wider
  - ▶ graphs and `lmer` analysis: longer.
- `pivot_wider`.

Making wide format

- pivot_wider needs: a column that is going to be split, and the column to make the values out of:

```
exercise.long %>% pivot_wider(names_from=time,
                              values_from=pulse) -> exercise.w
exercise.wide %>% sample_n(5) # random 5 rows
```

```
# A tibble: 5 x 6
     id diet     exertype min01 min15 min30
  <dbl> <chr>    <chr>    <dbl> <dbl> <dbl>
1    13 nonlowfat walking     90    92    93
2    28 lowfat   running    103   124   140
3    21 nonlowfat running     93    98   110
4    23 nonlowfat running     98   105    99
5    20 lowfat   walking    102   104   103
```

Setting up

- Make response variable from `min01`, `min15`, `min30`:

```
response <- with(exercise.wide, cbind(min01, min15, min30))
```

- Predict from `diet`, `exertype`, interaction using `lm`:

```
exercise.1 <- lm(response ~ diet * exertype,
  data = exercise.wide
)
```

… continued

- Run this through `Manova`:

```
times <- colnames(response)
times.df <- data.frame(times=factor(times))
exercise.2 <- Manova(exercise.1,
                     idata = times.df,
                     idesign = ~times)
```

Sphericity tests

```
summary(exercise.2)$sphericity.tests
```

```
                     Test statistic p-value
times                      0.92416 0.40372
diet:times                 0.92416 0.40372
exertype:times             0.92416 0.40372
diet:exertype:times        0.92416 0.40372
```

No problem with sphericity; go to univariate tests.

Univariate tests

```
summary(exercise.2)$univariate.tests
```

	Sum Sq	num Df	Error SS	den Df	F value
(Intercept)	894608	1	2085.2	24	10296.6595
diet	1262	1	2085.2	24	14.5238
exertype	8326	2	2085.2	24	47.9152
diet:exertype	816	2	2085.2	24	4.6945
times	2067	2	1563.6	48	31.7206
diet:times	193	2	1563.6	48	2.9597
exertype:times	2723	4	1563.6	48	20.9005
diet:exertype:times	614	4	1563.6	48	4.7095

	Pr(>F)	
(Intercept)	< 2.2e-16	***
diet	0.0008483	***
exertype	4.166e-09	***
diet:exertype	0.0190230	*
times	1.662e-09	***
diet:times	0.0613651	.
exertype:times	4.992e-10	***
diet:exertype:times	0.0027501	**

Comments

- The three-way interaction is significant
- the effect of diet on pulse rate over time is different for the different exercise types

Making some graphs

- Three-way interactions are difficult to understand. To make an attempt, look at some graphs.
- Plot time trace of pulse rates for each individual, joined by lines, and make *separate* plots for each diet-exertype combo.
- `facet_grid(diet~exertype)`: do a separate plot for each combination of diet and exercise type, with diets going down the page and exercise types going across. (Graphs are usually landscape, so have the factor `exertype` with more levels going across.)
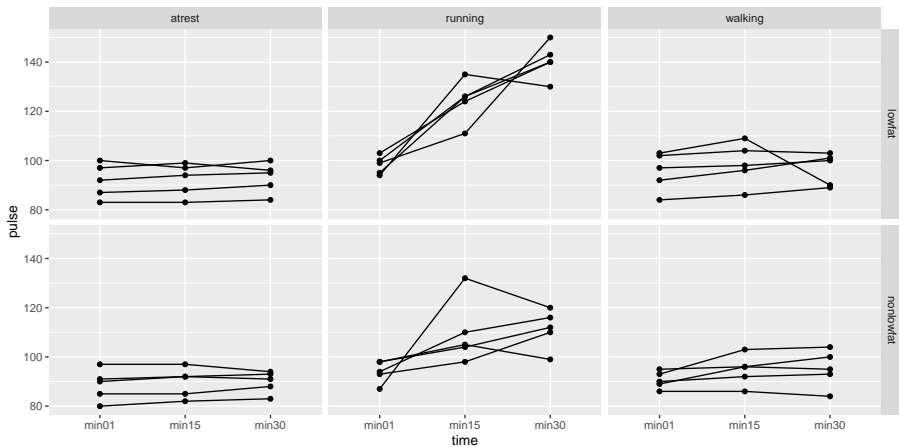
… continued

- ggplot again. Using *long* data frame:

```
g <- ggplot(exercise.long, aes(
  x = time, y = pulse,
  group = id
)) + geom_point() + geom_line() +
  facet_grid(diet ~ exertype)
```

The graph(s)

g

Comments on graphs

- At rest: no change in pulse rate over time
- Walking: not much change in pulse rates over time.
- Running: overall increase in pulse rate over time, but increase stronger for lowfat group.
- No consistent effect of:
  - ▶ diet over all exercise groups.
  - ▶ exercise type over both diet groups.
  - ▶ time over all diet-exercise type combos.

"Simple effects" of diet for the subjects who ran

- Looks as if there is only any substantial time effect for the runners. For them, does diet have an effect?
- Pull out only the runners from the wide data:

```
exercise.wide %>%
  filter(exertype == "running") -> runners.wide
```

… continued

- Create response variable and do MANOVA. Some of this looks like before, but I have different data now:

```
response <- with(runners.wide, cbind(min01, min15, min30))
runners.1 <- lm(response ~ diet, data = runners.wide)
times <- colnames(response)
times.df <- data.frame(times=factor(times))
runners.2 <- Manova(runners.1,
  idata = times.df,
  idesign = ~times
)
```

Sphericity tests

```
summary(runners.2)$sphericity.tests
```

```
           Test statistic p-value
times             0.81647  0.4918
diet:times        0.81647  0.4918
```

- No problem, look at univariate tests.

Univariate tests

```
summary(runners.2)$univariate.tests
```

```
            Sum Sq num Df Error SS den Df   F value     Pr(>F)
(Intercept) 383522      1    339.2      8 9045.3333 1.668e-13 ***
diet          1920      1    339.2      8   45.2830 0.0001482 ***
times         4714      2   1242.0     16   30.3644 3.575e-06 ***
diet:times     789      2   1242.0     16    5.0795 0.0195874 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Interaction still significant
  - dependence of pulse rate on time still different for the two diets

How is the effect of diet different over time?

- Table of means. Only I need long data for this:

```
runners.wide %>%
  pivot_longer(starts_with("min"),
               names_to = "time", values_to = "pulse") %>%
  group_by(time, diet) %>%
  summarize(
    mean = mean(pulse),
    sd = sd(pulse)
  ) -> summ
```

- Result of summarize is data frame, so can save it (and do more with it if needed).

Interaction plot

- We went to trouble of finding means by group, so making interaction plot is now mainly easy:

```
ggplot(summ, aes(x = time, y = mean, colour = diet,
                 group = diet)) + geom_point() + geom_line()
```

Comment on interaction plot

- The lines are not parallel, so there is interaction between diet and time for the runners.
- The effect of time on pulse rate is different for the two diets, even though all the subjects here were running.