

# Tidying data

## Tidying data

- Data rarely come to us as we want to use them.
- Before we can do analysis, typically have organizing to do.
- This is typical of ANOVA-type data, “wide format”:

```
pig feed1 feed2 feed3 feed4
 1 60.8 68.7 92.6 87.9
 2 57.0 67.7 92.1 84.2
 3 65.0 74.0 90.2 83.1
 4 58.6 66.3 96.5 85.7
 5 61.7 69.8 99.1 90.3
```

- 20 pigs randomly allocated to one of four feeds. At end of study, weight of each pig is recorded.
- Are any differences in mean weights among the feeds?
- Problem: want all weights in one column, with 2nd column labelling which feed. Untidy!

## Tidy and untidy data (Wickham)

- Data set easier to deal with if:
  - ▶ each observation is one row
  - ▶ each variable is one column
  - ▶ each type of observation unit is one table
- Data arranged this way called “tidy”; otherwise called “untidy”.
- For the pig data:
  - ▶ response variable is weight, but scattered over 4 columns, which are levels of a factor feed.
  - ▶ Want all the weights in one column, with a second column feed saying which feed that weight goes with.
  - ▶ Then we can run aov.

## Packages for this section

```
library(tidyverse)
```

## Reading in the pig data

```
my_url <- "http://ritsokiguess.site/datafiles/pigs1.txt"
pigs1 <- read_delim(my_url, " ")
pigs1
```

```
# A tibble: 5 x 5
  pig feed1 feed2 feed3 feed4
  <dbl> <dbl> <dbl> <dbl> <dbl>
1     1   60.8   68.7   92.6   87.9
2     2     57    67.7   92.1   84.2
3     3     65     74    90.2   83.1
4     4   58.6   66.3   96.5   85.7
5     5   61.7   69.8   99.1   90.3
```

## Making it longer

- We wanted all the weights in one column, labelled by which feed they went with.
- This is a very common reorganization, and the magic “verb” is `pivot_longer`:

```
pigs1 %>% pivot_longer(feed1:feed4, names_to="feed",  
                         values_to="weight") -> pigs2
```

# The long dataframe pigs2

```
# A tibble: 20 x 3
  pig feed  weight
  <dbl> <chr> <dbl>
1     1 feed1  60.8
2     1 feed2  68.7
3     1 feed3  92.6
4     1 feed4  87.9
5     2 feed1  57
6     2 feed2  67.7
7     2 feed3  92.1
8     2 feed4  84.2
9     3 feed1  65
10    3 feed2  74
11    3 feed3  90.2
12    3 feed4  83.1
13    4 feed1  58.6
14    4 feed2  66.3
15    4 feed3  96.5
16    4 feed4  85.7
17    5 feed1  61.7
18    5 feed2  69.8
19    5 feed3  99.1
20    5 feed4  90.3
```

## Alternatives

Any way of choosing the columns to pivot longer is good, eg:

```
pigs1 %>% pivot_longer(-pig, names_to="feed",
                           values_to="weight")
```

```
# A tibble: 20 x 3
  pig   feed  weight
  <dbl> <chr> <dbl>
1 1     feed1  60.8
2 1     feed2  68.7
3 1     feed3  92.6
4 1     feed4  87.9
5 2     feed1  57
6 2     feed2  67.7
7 2     feed3  92.1
8 2     feed4  84.2
9 3     feed1  65
10 3    feed2  74
```

## Comments

- pigs2 now in “long” format, ready for analysis.
- Anatomy of `pivot_longer`:
  - ▶ columns to combine
  - ▶ a name for column that will contain groups (“names”)
  - ▶ a name for column that will contain measurements (“values”)

## Identifying the pigs

- Values in `pig` identify pigs *within each group*: pig 1 is four different pigs!
- Create unique pig IDs by gluing pig number onto feed:

```
pigs2 %>% mutate(pig_id=str_c(feed, "_", pig)) -> pigs2
```

# The new pigs2

```
# A tibble: 20 x 4
  pig feed  weight pig_id
  <dbl> <chr>  <dbl> <chr>
1     1 feed1   60.8 feed1_1
2     1 feed2   68.7 feed2_1
3     1 feed3   92.6 feed3_1
4     1 feed4   87.9 feed4_1
5     2 feed1    57    feed1_2
6     2 feed2   67.7 feed2_2
7     2 feed3   92.1 feed3_2
8     2 feed4   84.2 feed4_2
9     3 feed1    65    feed1_3
10    3 feed2    74    feed2_3
11    3 feed3   90.2 feed3_3
12    3 feed4   83.1 feed4_3
13    4 feed1   58.6 feed1_4
14    4 feed2   66.3 feed2_4
15    4 feed3   96.5 feed3_4
16    4 feed4   85.7 feed4_4
17    5 feed1   61.7 feed1_5
18    5 feed2   69.8 feed2_5
19    5 feed3   99.1 feed3_5
20    5 feed4   90.3 feed4_5
```

## ...and finally, the analysis

- which is just what we saw before:

```
weight.1 <- aov(weight ~ feed, data = pigs2)
summary(weight.1)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
feed	3	3521	1173.5	119.1	3.72e-11	***
Residuals	16	158	9.8			
---						
Signif. codes:	0	'***'	0.001	'**'	0.01	'*'
					0.05	'.'
					0.1	' '
						1

- The mean weights of pigs on the different feeds are definitely not all equal.
- So we run Tukey to see which ones differ (over).

## Tukey

TukeyHSD(weight.1)

Tukey multiple comparisons of means  
95% family-wise confidence level

Fit: aov(formula = weight ~ feed, data = pigs2)

\$feed

	diff	lwr	upr	p	adj
feed2-feed1	8.68	3.001038	14.358962	0.0024000	
feed3-feed1	33.48	27.801038	39.158962	0.0000000	
feed4-feed1	25.62	19.941038	31.298962	0.0000000	
feed3-feed2	24.80	19.121038	30.478962	0.0000000	
feed4-feed2	16.94	11.261038	22.618962	0.0000013	
feed4-feed3	-7.86	-13.538962	-2.181038	0.0055599	

All of the feeds differ!

## Mean weights by feed

To find the best and worst, get mean weight by feed group. I borrowed an idea from earlier to put the means in descending order:

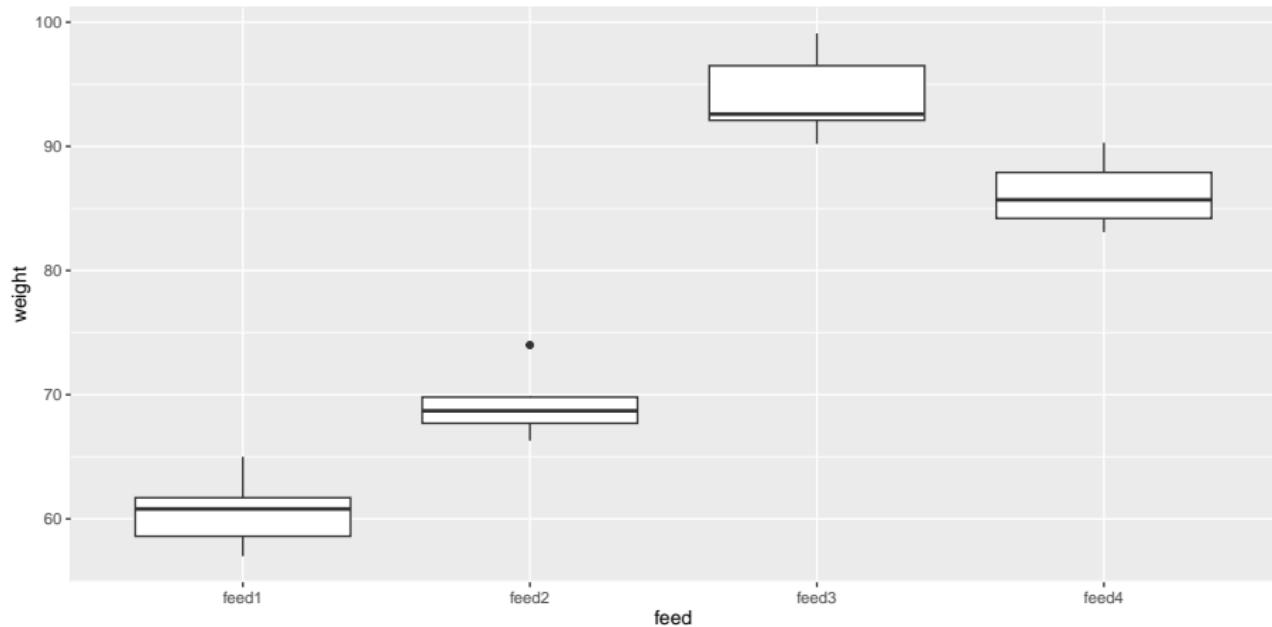
```
pigs2 %>%
  group_by(feed) %>%
  summarize(mean_weight = mean(weight))%>%
  arrange(desc(mean_weight))
```

```
# A tibble: 4 x 2
  feed   mean_weight
  <chr>     <dbl>
1 feed3      94.1
2 feed4      86.2
3 feed2      69.3
4 feed1      60.6
```

Feed 3 is best, feed 1 worst.

# Should we have any concerns about the ANOVA?

```
ggplot(pigs2, aes(x = feed, y = weight)) + geom_boxplot()
```



## Comments

- Feed 2 has an outlier
- But there are only 5 pigs in each group
- The conclusion is so clear that I am OK with this.

# Tuberculosis

- The World Health Organization keeps track of number of cases of various diseases, eg. tuberculosis.
- Some data:

```
my_url <- "http://ritsokiguess.site/datafiles/tb.csv"  
tb <- read_csv(my_url)
```

## The data (10 randomly chosen rows)

tb

```
# A tibble: 5,769 x 22
  iso2    year   m04   m514   m014   m1524   m2534   m3544   m4554   m5564   m65
  <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 AD     1989     NA     NA     NA     NA     NA     NA     NA     NA     NA
2 AD     1990     NA     NA     NA     NA     NA     NA     NA     NA     NA
3 AD     1991     NA     NA     NA     NA     NA     NA     NA     NA     NA
4 AD     1992     NA     NA     NA     NA     NA     NA     NA     NA     NA
5 AD     1993     NA     NA     NA     NA     NA     NA     NA     NA     NA
6 AD     1994     NA     NA     NA     NA     NA     NA     NA     NA     NA
7 AD     1996     NA     NA     0      0      0      4      1      0      0
8 AD     1997     NA     NA     0      0      1      2      2      1      6
9 AD     1998     NA     NA     0      0      0      1      0      0      0
10 AD    1999     NA     NA     0      0      0      1      1      0      0
# i 5,759 more rows
# i 11 more variables: mu <dbl>, f04 <dbl>, f514 <dbl>, f014 <dbl>,
#   f1524 <dbl>, f2534 <dbl>, f3544 <dbl>, f4554 <dbl>, f5564 <dbl>,
#   f65 <dbl>, fu <dbl>
```

## Many rows and columns

```
nrow(tb)
```

```
[1] 5769
```

```
ncol(tb)
```

```
[1] 22
```

## What we have

- Variables: country (abbreviated), year. Then number of cases for each gender and age group, eg. m1524 is males aged 15–24. Also mu and fu, where age is unknown.
- Lots of missings. Want to get rid of.
- Abbreviations [here](#).

```
tb %>%  
  pivot_longer(m04:fu, names_to = "genage",  
               values_to = "freq", values_drop_na = TRUE) -> t
```

- Code for `pivot_longer`:
  - ▶ columns to make longer
  - ▶ column to contain the names (categorical)
  - ▶ column to contain the values (quantitative)
  - ▶ drop missings in the values

## Results (some)

tb2

```
# A tibble: 35,750 x 4
  iso2    year genage freq
  <chr> <dbl> <chr>  <dbl>
1 AD      1996 m014      0
2 AD      1996 m1524     0
3 AD      1996 m2534     0
4 AD      1996 m3544     4
5 AD      1996 m4554     1
6 AD      1996 m5564     0
7 AD      1996 m65       0
8 AD      1996 f014      0
9 AD      1996 f1524     1
10 AD     1996 f2534     1
# i 35,740 more rows
```

# Separating

- 4 columns, but 5 variables, since genage contains both gender and age group. Split that up using separate.
- `separate` needs to know:
  - ▶ what to separate (no quotes needed),
  - ▶ how to split, and what to separate into (here you do need quotes):

```
tb2 %>%  
  separate_wider_position(genage,  
    widths = c("gender" = 1, "age" = 4),  
    too_few = "align_start") -> tb3
```

## Tidied tuberculosis data (some)

tb3

```
# A tibble: 35,750 x 5
  iso2    year gender age     freq
  <chr> <dbl> <chr> <chr> <dbl>
1 AD      1996 m     014      0
2 AD      1996 m     1524     0
3 AD      1996 m     2534     0
4 AD      1996 m     3544     4
5 AD      1996 m     4554     1
6 AD      1996 m     5564     0
7 AD      1996 m     65       0
8 AD      1996 f     014      0
9 AD      1996 f     1524     1
10 AD     1996 f     2534     1
# i 35,740 more rows
```

## In practice...

- instead of doing the pipe one step at a time, you *debug* it one step at a time, and when you have each step working, you use that step's output as input to the next step, thus:

```
tb %>%  
  pivot_longer(m04:fu, names_to = "genage",  
              values_to = "freq", values_drop_na = TRUE) %>%  
  separate_wider_position(genage,  
                          widths = c("gender" = 1, "age" = 4),  
                          too_few = "align_start") -> tb_tidy
```

- When you have it working, save the final result (for further work).

## Comments

- You can split the R code over as many lines as you like, as long as each line is incomplete, so that R knows more is to come.
- I like to put the pipe symbol on the end of the line.
- Sometimes one function call gets very long, in which case you can separate at commas.

## Total tuberculosis cases by year (some of the years)

```
tb_tidy %>%
  filter(between(year, 1991, 1998)) %>%
  group_by(year) %>% summarize(total=sum(freq))
```

```
# A tibble: 8 x 2
```

	year	total
	<dbl>	<dbl>
1	1991	544
2	1992	512
3	1993	492
4	1994	750
5	1995	513971
6	1996	635705
7	1997	733204
8	1998	840389

- Something very interesting happened between 1994 and 1995.

## To find out what

- try counting up total cases by country:

```
tb_tidy %>% group_by(iso2) %>%
  summarize(total=sum(freq)) %>%
  arrange(desc(total))
```

```
# A tibble: 213 x 2
```

```
  iso2      total
  <chr>    <dbl>
1 CN      4065174
2 IN      3966169
3 ID      1129015
4 ZA      900349
5 BD      758008
6 VN      709695
7 CD      603095
8 PH      490040
9 BR      440609
10 KE     431523
# i 203 more rows
```

## What years do I have for China?

China started recording in 1995, which is at least part of the problem:

```
tb_tidy %>% filter(iso2 == "CN") %>%
  group_by(year) %>%
  summarize(total = sum(freq))
```

```
# A tibble: 14 x 2
  year   total
  <dbl>   <dbl>
1 1995 131194
2 1996 168270
3 1997 195895
4 1998 214404
5 1999 212258
6 2000 213766
7 2001 212766
8 2002 194972
9 2003 267280
10 2004 384886
11 2005 472719
```

## First year of recording by country?

- A lot of countries started recording in about 1995, in fact:

```
tb_tidy %>% group_by(iso2) %>%
  summarize(first_year=min(year)) %>%
  count(first_year)
```

```
# A tibble: 14 x 2
  first_year     n
  <dbl> <int>
1 1980          2
2 1994          2
3 1995         130
4 1996          31
5 1997          17
6 1998           6
7 1999          10
8 2000           4
9 2001           1
10 2002          3
11 2003          2
```

## Comment

- So the reason for the big jump in cases is that so many countries started recording then, not that there really were more cases.

## Some Toronto weather data

```
my_url <- "http://ritsokiguess.site/STAC32/toronto_weather.csv"
weather <- read_csv(my_url)
weather

# A tibble: 24 x 35
  station Year Month element   d01   d02   d03   d04   d05   d06   d07
  <chr>    <dbl> <chr> <chr>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 TORONT~  2018  01   tmax     -7.9  -7.1  -5.3  -7.7 -14.7 -15.4  -1
2 TORONT~  2018  01   tmin    -18.6 -12.5 -11.2 -19.7 -20.6 -22.3 -17.5
3 TORONT~  2018  02   tmax      5.6  -8.6   0.4   1.8  -6.6  -3.2  -4.1
4 TORONT~  2018  02   tmin    -8.9  -15   -9.7  -8.8  -12   -8.2  -8.7
5 TORONT~  2018  03   tmax      NA    NA    NA    NA    NA    NA    3.1
6 TORONT~  2018  03   tmin      NA   -0.5  NA   -3.1  NA   -1.4  0.4
7 TORONT~  2018  04   tmax      4.5   6.5   5    5.7   2.9   5.4   2
8 TORONT~  2018  04   tmin    -2.6  -1.2   2.4  -3.2  -3.9  -2.6  -4.4
9 TORONT~  2018  05   tmax    23.5  26.3  23    24  24.1  17.4  15.9
10 TORONT~ 2018  05   tmin     8.5  14.4  11.4   9.2   8.5  13.3  10.6
# i 14 more rows
# i 24 more variables: d08 <dbl>, d09 <dbl>, d10 <dbl>, d11 <dbl>,
#   d12 <dbl>, d13 <dbl>, d14 <dbl>, d15 <dbl>, d16 <dbl>, d17 <dbl>,
#   d18 <dbl>, d19 <dbl>, d20 <dbl>, d21 <dbl>, d22 <dbl>, d23 <dbl>,
#   d24 <dbl>, d25 <dbl>, d26 <dbl>, d27 <dbl>, d28 <dbl>, d29 <dbl>, d30 <dbl>
```

## The columns

- Daily weather records for “Toronto City” weather station in 2018:
  - ▶ station: identifier for this weather station (always same here)
  - ▶ Year, Month
  - ▶ element: whether temperature given was daily max or daily min
  - ▶ d01, d02,... d31: day of the month from 1st to 31st.

## Off we go

Numbers in data frame all temperatures (for different days of the month), so first step is

```
weather %>%  
  pivot_longer(d01:d31, names_to="day",  
               values_to="temperature",  
               values_drop_na = TRUE)
```

```
# A tibble: 703 x 6  
  station      Year Month element day   temperature  
  <chr>        <dbl> <chr>  <chr>  <chr>       <dbl>  
1 TORONTO CITY 2018  01    tmax   d01    -7.9  
2 TORONTO CITY 2018  01    tmax   d02    -7.1  
3 TORONTO CITY 2018  01    tmax   d03    -5.3  
4 TORONTO CITY 2018  01    tmax   d04    -7.7  
5 TORONTO CITY 2018  01    tmax   d05   -14.7  
6 TORONTO CITY 2018  01    tmax   d06   -15.4  
7 TORONTO CITY 2018  01    tmax   d07    -1  
8 TORONTO CITY 2018  01    tmax   d08     3  
9 TORONTO CITY 2018  01    tmax   d09    1.6
```

## Element

- Column element contains names of two different variables, that should each be in separate column.
- Distinct from eg. `m1524` in tuberculosis data, that contained levels of two different factors, handled by separate.
- Untangling names of variables handled by `pivot_wider`.

## Handling element

```
weather %>%
  pivot_longer(d01:d31, names_to="day",
               values_to="temperature",
               values_drop_na = TRUE) %>%
  pivot_wider(names_from=element,
              values_from=temperature)
```

```
# A tibble: 355 x 6
  station      Year Month day     tmax     tmin
  <chr>       <dbl> <chr> <chr>   <dbl>   <dbl>
1 TORONTO CITY 2018  01    d01    -7.9   -18.6
2 TORONTO CITY 2018  01    d02    -7.1   -12.5
3 TORONTO CITY 2018  01    d03    -5.3   -11.2
4 TORONTO CITY 2018  01    d04    -7.7   -19.7
5 TORONTO CITY 2018  01    d05   -14.7   -20.6
6 TORONTO CITY 2018  01    d06   -15.4   -22.3
7 TORONTO CITY 2018  01    d07    -1     -17.5
```

## Further improvements 1/2

- We have tidy data now, but can improve things further.
- `mutate` creates new columns from old (or assign back to change a variable).
- Would like numerical dates. `separate` works, or pull out number as below.
- `select` keeps columns (or drops, with minus). Station name has no value to us.

## Further improvements 2/2

```
weather %>%
  pivot_longer(d01:d31, names_to="day",
               values_to="temperature", values_drop_na = TRUE) %>%
  pivot_wider(names_from=element, values_from=temperature) %>%
  mutate(Day = parse_number(day)) %>%
  select(-station)
```

```
# A tibble: 355 x 6
  Year Month day     tmax    tmin   Day
  <dbl> <chr> <chr> <dbl> <dbl> <dbl>
1 2018 01   d01    -7.9  -18.6    1
2 2018 01   d02    -7.1  -12.5    2
3 2018 01   d03    -5.3  -11.2    3
4 2018 01   d04    -7.7  -19.7    4
5 2018 01   d05   -14.7  -20.6    5
6 2018 01   d06   -15.4  -22.3    6
7 2018 01   d07     -1   -17.5    7
8 2018 01   d08      3   -1.7     8
9 2018 01   d09     1.6  -0.6     9
10 2018 01  d10     5.9  -1.3    10
```

## Final step(s)

- Make year-month-day into proper date.
- Keep only date, tmax, tmin:

```
weather %>%
  pivot_longer(d01:d31, names_to="day",
               values_to="temperature", values_drop_na = T) %>%
  pivot_wider(names_from=element, values_from=temperature) %>%
  mutate(Day = parse_number(day)) %>%
  select(-station) %>%
  unite(datestr, c(Year, Month, Day), sep = "-") %>%
  mutate(date = as.Date(datestr)) %>%
  select(date, tmax, tmin) -> weather_tidy
```

## Our tidy data frame

```
weather_tidy
```

```
# A tibble: 355 x 3
  date       tmax   tmin
  <date>     <dbl>  <dbl>
1 2018-01-01 -7.9 -18.6
2 2018-01-02 -7.1 -12.5
3 2018-01-03 -5.3 -11.2
4 2018-01-04 -7.7 -19.7
5 2018-01-05 -14.7 -20.6
6 2018-01-06 -15.4 -22.3
7 2018-01-07 -1     -17.5
8 2018-01-08  3     -1.7
9 2018-01-09  1.6   -0.6
10 2018-01-10  5.9   -1.3
# i 345 more rows
```

## Plotting the temperatures

- Plot temperature against date joined by lines, but with separate lines for max and min. ggplot requires something like

```
ggplot(..., aes(x = date, y = temperature)) + geom_point() +  
  geom_line()
```

only we have two temperatures, one a max and one a min, that we want to keep separate.

- The trick: combine `tmax` and `tmin` together into one column, keeping track of what kind of temp they are. (This actually same format as untidy `weather`.) Are making `weather_tidy` untidy for purposes of drawing graph only.
- Then can do something like

```
ggplot(d, aes(x = date, y = temperature, colour = maxmin))  
  + geom_point() + geom_line()
```

to distinguish max and min on graph.

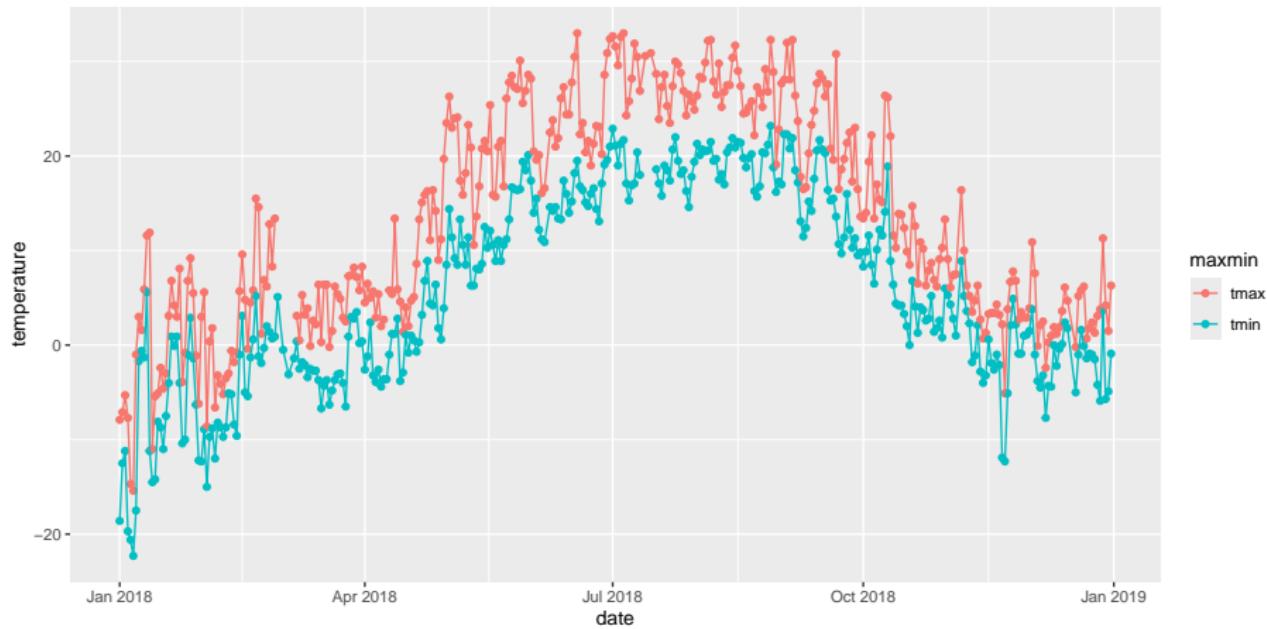
## Setting up plot

- Since we only need data frame for plot, we can do the column-creation and plot in a pipeline.
- For a ggplot in a pipeline, the initial data frame is omitted, because it is whatever came out of the previous step.
- To make those “one column”s: pivot\_longer. I save the graph to show overleaf:

```
weather_tidy %>%
  pivot_longer(tmax:tmin, names_to="maxmin",
              values_to="temperature") %>%
  ggplot(aes(x = date, y = temperature, colour = maxmin)) +
  geom_point() + geom_line() -> g
```

# The plot

gg



## Summary of tidying “verbs”

---

Verb	Purpose
<code>pivot_longer</code>	Combine columns that measure same thing into one
<code>pivot_wider</code>	Take column that measures one thing under different conditions and put into multiple columns
<code>separate</code>	Turn a column that encodes several variables into several columns
<code>unite</code>	Combine several (related) variables into one “combination” variable

---

`pivot_longer` and `pivot_wider` are opposites; `separate` and `unite` are opposites.