Regression with categorical variables

Packages for this section

```
library(tidyverse)
library(broom)
```

The pigs revisited

- Recall pig feed data, after we tidied it:

```
my_url <- "http://ritsokiguess.site/datafiles/pigs2.txt"
pigs <- read_delim(my_url, " ")
pigs
```

```
# A tibble: 20 x 3
     pig feed  weight
   <dbl> <chr>  <dbl>
 1     1 feed1   60.8
 2     2 feed1   57
 3     3 feed1   65
 4     4 feed1   58.6
 5     5 feed1   61.7
 6     1 feed2   68.7
 7     2 feed2   67.7
 8     3 feed2   74
 9     4 feed2   66.3
```

Summaries

```r
pigs %>%
  group_by(feed) %>%
  summarize(n = n(), mean_wt = mean(weight),
            sd_wt = sd(weight))
```

```
# A tibble: 4 x 4
  feed      n mean_wt sd_wt
  <chr> <int>   <dbl> <dbl>
1 feed1     5    60.6  3.06
2 feed2     5    69.3  2.93
3 feed3     5    94.1  3.61
4 feed4     5    86.2  2.90
```

Running through `aov` and `lm`

- What happens if we run this through `lm` rather than `aov`?
- Recall `aov` first:

```
pigs.1 <- aov(weight ~ feed, data = pigs)
summary(pigs.1)
```

```
            Df Sum Sq Mean Sq F value   Pr(>F)
feed         3   3521  1173.5   119.1 3.72e-11 ***
Residuals   16    158     9.9
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

and now lm

```
pigs.2 <- lm(weight ~ feed, data = pigs)
summary(pigs.2)
```

```
Call:
lm(formula = weight ~ feed, data = pigs)

Residuals:
   Min     1Q Median     3Q    Max
-3.900 -2.025 -0.570  1.845  5.000

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)   60.620      1.404  43.190  < 2e-16 ***
feedfeed2      8.680      1.985   4.373 0.000473 ***
feedfeed3     33.480      1.985  16.867 1.30e-11 ***
feedfeed4     25.620      1.985  12.907 7.11e-10 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.138 on 16 degrees of freedom
```

Understanding those slopes

- Get one slope for each category of categorical variable feed, except for first.
- feed1 treated as "baseline", others measured relative to that.
- Thus prediction for feed 1 is intercept, 60.62 (mean weight for feed 1).
- Prediction for feed 2 is $60.62 + 8.68 = 69.30$ (mean weight for feed 2).
- Or, mean weight for feed 2 is 8.68 bigger than for feed 1.
- Mean weight for feed 3 is 33.48 bigger than for feed 1.
- Slopes can be negative, if mean for a feed had been smaller than for feed 1.

Reproducing the ANOVA

- Pass the fitted model object into anova:

```
anova(pigs.2)
```

```
Analysis of Variance Table

Response: weight
          Df Sum Sq Mean Sq F value   Pr(>F)
feed       3 3520.5 1173.51  119.14 3.72e-11 ***
Residuals 16  157.6    9.85
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Same as before.
- But no Tukey this way:

```
TukeyHSD(pigs.2)
```

```
Error in UseMethod("TukeyHSD"): no applicable method for 'TukeyHSD' applied
```

The crickets

- Male crickets rub their wings together to produce a chirping sound.
- Rate of chirping, called "pulse rate", depends on species and possibly on temperature.
- Sample of crickets of two species' pulse rates measured; temperature also recorded.
- Does pulse rate differ for species, especially when temperature accounted for?

The crickets data

Read the data:

```
my_url <- "http://ritsokiguess.site/datafiles/crickets2.csv"
crickets <- read_csv(my_url)
crickets %>% slice_sample(n = 10)
```

```
# A tibble: 10 x 3
   species        temperature pulse_rate
   <chr>                <dbl>      <dbl>
 1 exclamationis         26.2       86.6
 2 exclamationis         26.2       85.8
 3 exclamationis         30.4      102.
 4 niveus                26.5       77
 5 niveus                26.5       76.1
 6 niveus                26.5       77.7
 7 niveus                21         58.5
 8 niveus                18.3       47.6
 9 niveus                22.1       60.7
```

Fit model with `lm`

```
crickets.1 <- lm(pulse_rate ~ temperature + species,
                 data = crickets)
summary(crickets.1)
```

```
Call:
lm(formula = pulse_rate ~ temperature + species, data = cricke

Residuals:
    Min      1Q  Median      3Q     Max
-3.0128 -1.1296 -0.3912  0.9650  3.7800

Coefficients:
               Estimate Std. Error t value Pr(>|t|)
(Intercept)    -7.21091    2.55094  -2.827  0.00858 **
temperature     3.60275    0.09729  37.032  < 2e-16 ***
speciesniveus -10.06529    0.73526 -13.689 6.27e-14 ***
```

The summary

```
summary(crickets.1)
```

```
Call:
lm(formula = pulse_rate ~ temperature + species, data = cricke

Residuals:
   Min      1Q  Median      3Q     Max
-3.0128 -1.1296 -0.3912  0.9650  3.7800

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   -7.21091    2.55094  -2.827  0.00858 **
temperature    3.60275    0.09729  37.032  < 2e-16 ***
speciesniveus -10.06529    0.73526 -13.689 6.27e-14 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Conclusions

- Slope for temperature says that increasing temperature by 1 degree increases pulse rate by 3.6 (same for both species)
- Slope for `speciesniveus` says that pulse rate for `niveus` about 10 lower than that for `exclamationis` at same temperature (latter species is baseline).
- R-squared of almost 0.99 is very high, so that the prediction of pulse rate from species and temperature is very good.

To end with a graph

- Two quantitative variables and one categorical: scatterplot with categories distinguished by colour.
- This graph seems to need a title, which I define first.

```
t1 <- "Pulse rate against temperature for two species of crick
t2 <- "Temperature in degrees Celsius"
ggplot(crickets, aes(x = temperature, y = pulse_rate,
  colour = species)) +
  geom_point() + geom_smooth(method = "lm", se = FALSE) -> g
```

The graph

g